

BUSINESS UNDERSTANDING.

Customer churn refers to when a customer stops using a company's services. In the telecommunications industry, this is a critical issue as it directly impacts the company's revenue and growth. For SyriaTeL, it is important to understand the churn rate to be able to manage the customer base to ensure long-term profitability.

Business Problem Statement

SyriaTel is facing a level of customer turnover, resulting in substantial revenue decline. This study seeks to address this challenge by analysing the different factors that contribute to this challenge. Ultimately, the goal is to forecast which customers are likely to discontinue their service and deploy effective retention measures to mitigate this rate. We will be building a classification system and to evaluate the performance of the classifier we will use classification metrics including accuracy, recall, f1score and precision.

Objectives

1. Create different models to discover the patterns which lead to a high rate of customer churn.
2. Develop targeted retention strategies to enhance customer loyalty.

DATA UNDERSTANDING.

We first import the Churn dataset that we will be using for this project.

```
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
from scipy.stats import zscore
```

```
data = pd.read_csv(r"C:\Users\Administrator\Downloads\
bigml_59c28831336c6604c800002a.csv")
data.head()
```

	state	account length	area	code	phone number	international	plan	\
0	KS	128		415	382-4657		no	
1	OH	107		415	371-7191		no	
2	NJ	137		415	358-1921		no	
3	OH	84		408	375-9999		yes	
4	OK	75		415	330-6626		yes	

	voice mail plan	number vmail messages	total day minutes	total day calls	\
--	-----------------	-----------------------	-------------------	-----------------	---

0	yes	25	265.1
110			
1	yes	26	161.6
123			
2	no	0	243.4
114			
3	no	0	299.4
71			
4	no	0	166.7
113			
total day charge ... total eve calls total eve charge \			
0	45.07 ...	99	16.78
1	27.47 ...	103	16.62
2	41.38 ...	110	10.30
3	50.90 ...	88	5.26
4	28.34 ...	122	12.61
total night minutes total night calls total night charge \			
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41
total intl minutes total intl calls total intl charge \			
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73
customer service calls churn			
0	1	False	
1	1	False	
2	0	False	
3	2	False	
4	3	False	
[5 rows x 21 columns]			

COLUMN NAMES AND DESCRIPTION OF THE DATASET.

The dataset contains data on the customers of a SyriaTel company. Each row represents a customer and the columns contain customer's attributes which are described in the following:

- state: the state the user lives in
- account length: the number of days the user has this account
- area code: the code of the area the user lives in
- phone number: the phone number of the user

- international plan: true if the user has the international plan.
- voice mail plan: true if the user has the voice mail plan.
- number vmail messages: the number of voice mail messages the user has sent.
- total day minutes: total number of minutes the user has been in calls during the day.
- total day calls: total number of calls the user has done during the day.
- total day charge: total amount of money the user was charged by SyriaTel company for calls during the day.
- total eve minutes: total number of minutes the user has been in calls during the evening.
- total eve calls: total number of calls the user has done during the evening.
- total eve charge: total amount of money the user was charged by SyriaTel company for calls during the evening.
- total night minutes: total number of minutes the user has been in calls during the night.
- total night calls: total number of calls the user has done during the night.
- total night charge: total amount of money the user was charged by SyriaTel company for calls during the night.
- total intl minutes: total number of minutes the user has been in international calls.
- total intl calls: total number of international calls the user has done.
- total intl charge: total amount of money the user was charged by the Telecom company for international calls.
- customer service calls: number of customer service calls the user has done.
- churn: true if the user terminated the contract, or false the user is still using the company's services.

```
data.tail()
```

	state	account length	area code	phone number	international plan
3328	AZ	192	415	414-4276	no
3329	WV	68	415	370-3271	no
3330	RI	28	510	328-8230	no
3331	CT	184	510	364-6381	yes
3332	TN	74	415	400-4344	no

	voice mail plan	number vmail messages	total day minutes
3328	yes	36	156.2
3329	no	0	231.1
3330	no	0	180.8
3331	no	0	213.8
3332	yes	25	234.4

	total day calls	total day charge	...	total eve calls
3328	77	26.55	...	126
3329	57	39.29	...	55

3330	109	30.74	...	58
3331	105	36.35	...	84
3332	113	39.85	...	82

	total eve charge	total night minutes	total night calls	\
3328	18.32	279.1	83	
3329	13.04	191.3	123	
3330	24.55	191.9	91	
3331	13.57	139.2	137	
3332	22.60	241.4	77	

	total night charge	total intl minutes	total intl calls	\
3328	12.56	9.9	6	
3329	8.61	9.6	4	
3330	8.64	14.1	6	
3331	6.26	5.0	10	
3332	10.86	13.7	4	

	total intl charge	customer service calls	churn
3328	2.67	2	False
3329	2.59	3	False
3330	3.81	2	False
3331	1.35	2	False
3332	3.70	0	False

[5 rows x 21 columns]

data.info()

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 3333 entries, 0 to 3332

Data columns (total 21 columns):

#	Column	Non-Null Count	Dtype
---	-----	-----	-----
0	state	3333 non-null	object
1	account length	3333 non-null	int64
2	area code	3333 non-null	int64
3	phone number	3333 non-null	object
4	international plan	3333 non-null	object
5	voice mail plan	3333 non-null	object
6	number vmail messages	3333 non-null	int64
7	total day minutes	3333 non-null	float64
8	total day calls	3333 non-null	int64
9	total day charge	3333 non-null	float64
10	total eve minutes	3333 non-null	float64
11	total eve calls	3333 non-null	int64
12	total eve charge	3333 non-null	float64
13	total night minutes	3333 non-null	float64
14	total night calls	3333 non-null	int64
15	total night charge	3333 non-null	float64

```

16 total intl minutes      3333 non-null float64
17 total intl calls       3333 non-null int64
18 total intl charge      3333 non-null float64
19 customer service calls 3333 non-null int64
20 churn                  3333 non-null bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

data.shape

(3333, 21)

```

The dataset has 3333 rows and 21 columns .It comprised of different data types i.e integer, floats, boolean and strings.

DATA CLEANING

Data cleaning involves the process of identifying and resolving issues related to the quality of the dataset. Its primary objective is to ensure that the data is accurate, consistent, and devoid of errors.

#creating a copy of the dataset first.

```
df = data.copy()
df.head()
```

	state	account length	area code	phone number	international plan	\
0	KS	128	415	382-4657	no	
1	OH	107	415	371-7191	no	
2	NJ	137	415	358-1921	no	
3	OH	84	408	375-9999	yes	
4	OK	75	415	330-6626	yes	

	voice mail plan	number vmail messages	total day minutes	total day calls	\
0	yes	25	265.1		
110					
1	yes	26	161.6		
123					
2	no	0	243.4		
114					
3	no	0	299.4		
71					
4	no	0	166.7		
113					

	total day charge	...	total eve calls	total eve charge	\
0	45.07	...	99	16.78	
1	27.47	...	103	16.62	

2	41.38	...	110	10.30
3	50.90	...	88	5.26
4	28.34	...	122	12.61

	total night minutes	total night calls	total night charge \
0	244.7	91	11.01
1	254.4	103	11.45
2	162.6	104	7.32
3	196.9	89	8.86
4	186.9	121	8.41

	total intl minutes	total intl calls	total intl charge \
0	10.0	3	2.70
1	13.7	3	3.70
2	12.2	5	3.29
3	6.6	7	1.78
4	10.1	3	2.73

	customer service calls	churn
0	1	False
1	1	False
2	0	False
3	2	False
4	3	False

[5 rows x 21 columns]

#Checking for missing values.
df.isnull().sum()

```
state 0
account length 0
area code 0
phone number 0
international plan 0
voice mail plan 0
number vmail messages 0
total day minutes 0
total day calls 0
total day charge 0
total eve minutes 0
total eve calls 0
total eve charge 0
total night minutes 0
total night calls 0
total night charge 0
total intl minutes 0
total intl calls 0
total intl charge 0
customer service calls 0
```

```
churn          0
dtype: int64
```

The dataset has no missing values.

We will then change the column names and make them into titles.

```
#making the columns into titles and changing
#some of the column names to full column names.

df.rename(columns={
    'total eve minutes': 'total evening minutes',
    'number vmail messages': 'number of voicemail messages',
    'total eve charge': 'total evening charge',
    'total eve calls' : 'total evening calls',
    'total intl calls' : ' total international calls',
    'total intl charge' : 'total international charge',
    'total intl minutes' : 'total international minutes'
}, inplace=True)

df.columns = df.columns.str.title()

df.columns
Index(['State', 'Account Length', 'Area Code', 'Phone Number',
       'International Plan', 'Voice Mail Plan', 'Number Of Voicemail
Messages',
       'Total Day Minutes', 'Total Day Calls', 'Total Day Charge',
       'Total Evening Minutes', 'Total Evening Calls', 'Total Evening
Charge',
       'Total Night Minutes', 'Total Night Calls', 'Total Night
Charge',
       'Total International Minutes', ' Total International Calls',
       'Total International Charge', 'Customer Service Calls',
       'Churn'],
      dtype='object')
```

Checking for duplicate values.

```
df.duplicated().any
<bound method Series.any of 0      False
1      False
2      False
3      False
4      False
...
3328   False
3329   False
3330   False
```

```
3331     False
3332     False
Length: 3333, dtype: bool>
```

There are no duplicates in this dataset

We will then check for placeholders in the dataset.

These denote missing values in the dataset. Common placeholders used are "", 'N/A', 'NULL', '?', 'nan', 'unknown'. Identifying placeholders will help us to improve model performance and get more accurate analysis.

```
# Define a list of potential placeholder values
common_placeholders = ["", "NA", "N/A", "nan", "none", "null", "?",
                        "unknown", "missing"]

# Loop through each column and check for potential placeholders
found_placeholder = False
for column in df.columns:
    unique_values = df[column].unique()
    for value in unique_values:
        if pd.isna(value) or (isinstance(value, str) and
                               value.strip().lower() in common_placeholders):
            count = (df[column] == value).sum()
            print(f"Column '{column}': Found {count} occurrences of
potential placeholder '{value}'")
            found_placeholder = True

if not found_placeholder:
    print("No placeholders found in the DataFrame.")
```

No placeholders found in the DataFrame.

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 3333 entries, 0 to 3332
```

```
Data columns (total 21 columns):
```

#	Column	Non-Null Count	Dtype
0	State	3333 non-null	object
1	Account Length	3333 non-null	int64
2	Area Code	3333 non-null	int64
3	Phone Number	3333 non-null	object
4	International Plan	3333 non-null	object
5	Voice Mail Plan	3333 non-null	object
6	Number Of Voicemail Messages	3333 non-null	int64
7	Total Day Minutes	3333 non-null	float64
8	Total Day Calls	3333 non-null	int64
9	Total Day Charge	3333 non-null	float64


```

10 Total Evening Minutes      3333 non-null float64
11 Total Evening Calls        3333 non-null int64
12 Total Evening Charge       3333 non-null float64
13 Total Night Minutes        3333 non-null float64
14 Total Night Calls          3333 non-null int64
15 Total Night Charge         3333 non-null float64
16 Total International Minutes 3333 non-null float64
17 Total International Calls   3333 non-null int64
18 Total International Charge  3333 non-null float64
19 Customer Service Calls     3333 non-null int64
20 Churn                      3333 non-null bool
dtypes: bool(1), float64(8), int64(8), object(4)
memory usage: 524.2+ KB

```

#Drop the phone number column.

```
df = df.drop(['Phone Number'],axis=1)
```

We drop the phone number column since it is a unique identifier for each customer , it does not provide any meaningful information about a customer's characteristics.

We will change the Area Code column datatype from integer to object. It should be considered as a categorical column since it is a nominal variable. Also convert the churn column to an integer. And also change the datatype of churn from boolean to int.

```
df['Area Code'] = df['Area Code'].astype(str)
df['Churn'] = df['Churn'].astype(int)
```

Let us check for outliers on numeric columns. We will visualize the data to see its distribution. We will then handle the outliers as they will significantly affect the models' performance.

```
numeric_columns = df.select_dtypes(include=['float64', 'int64',
'int32'])
```

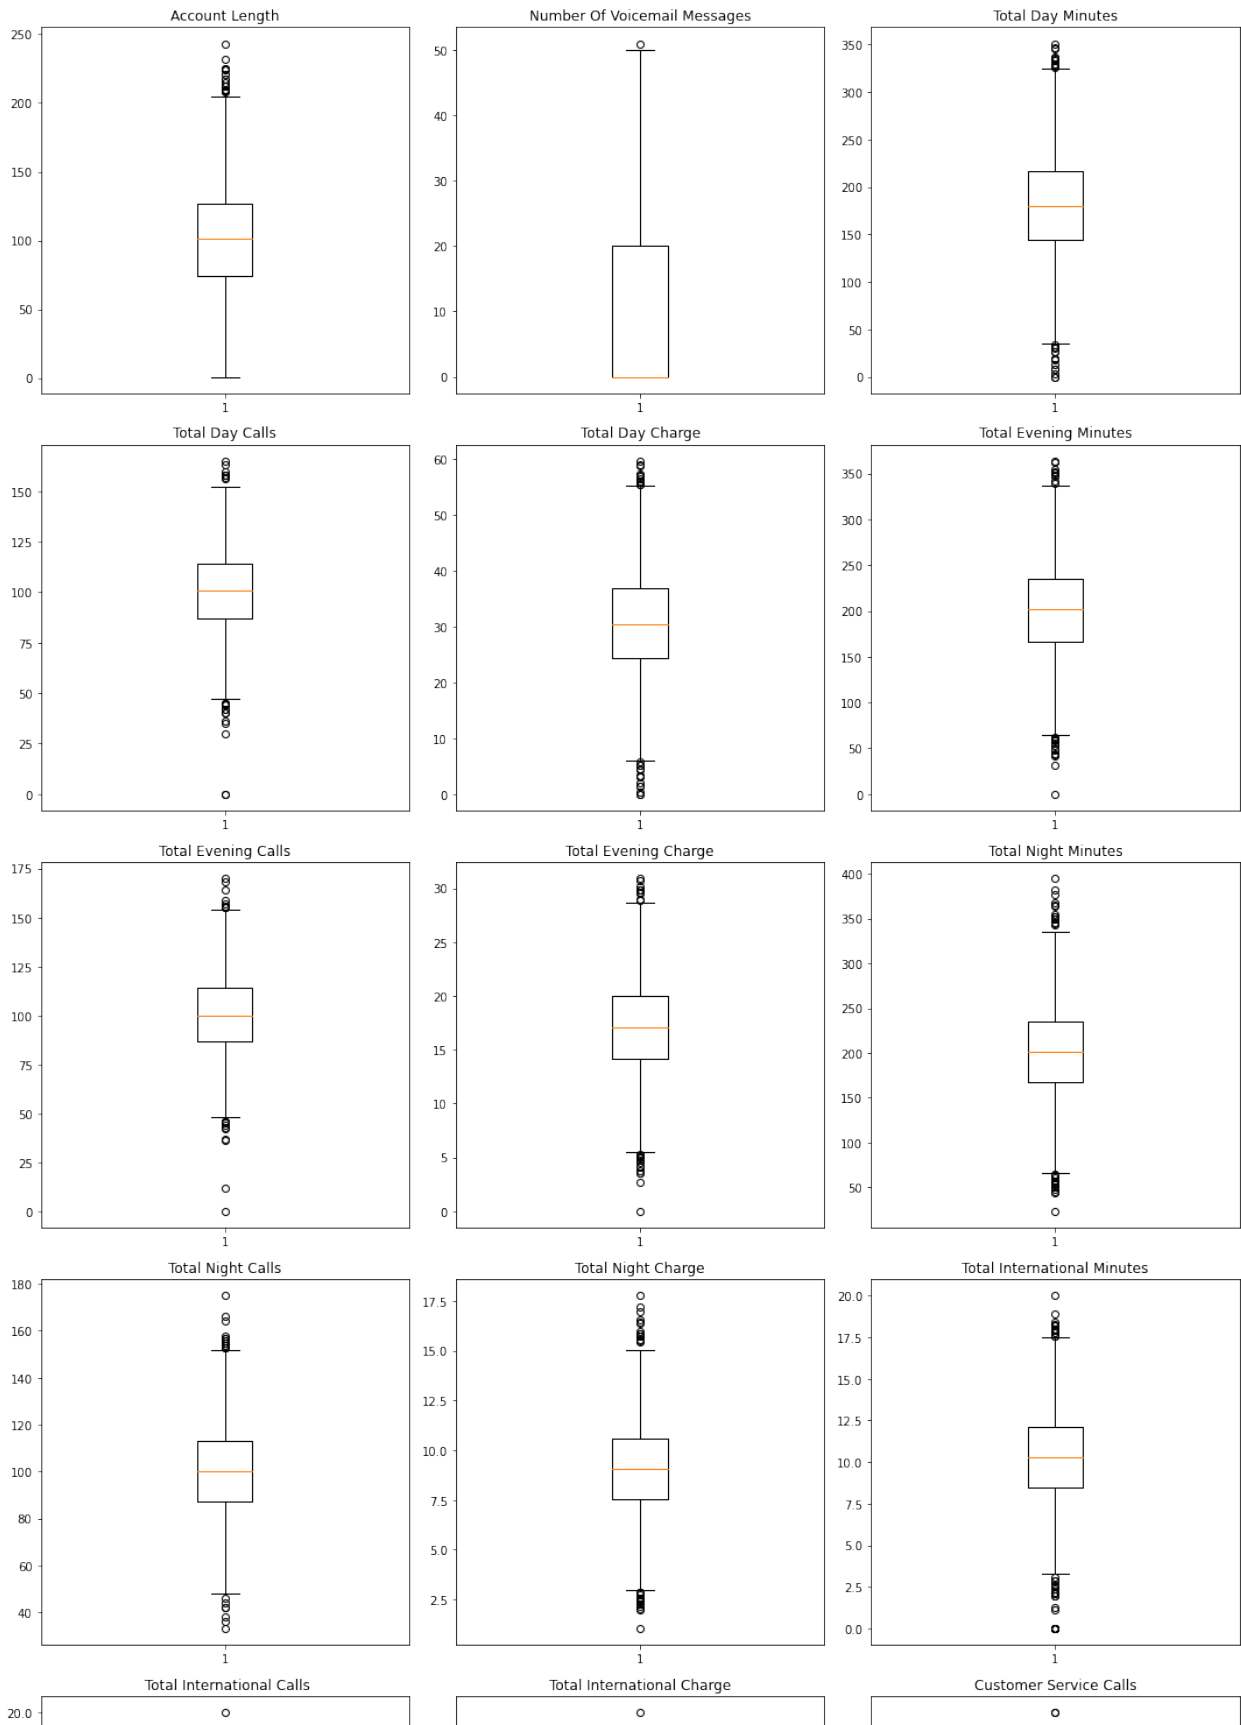
Plot box plots for each numeric column

```

num_cols = len(numeric_columns.columns)
cols_per_row = 3
num_rows = (num_cols - 1) // cols_per_row + 1

plt.figure(figsize=(15, 5 * num_rows))
for i, col in enumerate(numeric_columns.columns):
    plt.subplot(num_rows, cols_per_row, i+1)
    plt.boxplot(numeric_columns[col])
    plt.title(col)
plt.tight_layout()
plt.show()

```



We will use the z-score method to identify rows which have outliers.

```
z_scores = np.abs(zscore(df.select_dtypes(include=[np.number])))  
  
# Set threshold for Z-scores  
threshold = 3  
  
# Identify rows with any Z-score above the threshold  
outliers = (z_scores > threshold).any(axis=1)  
  
# Print the identified outliers  
print("Outliers detected using Z-score:")  
print(df[outliers])
```

Outliers detected using Z-score:

	State	Account Length	Area Code	International Plan	Voice Mail
Plan \					
22	AZ	130	415	no	
no					
32	LA	172	408	no	
no					
41	MD	135	408	yes	
yes					
58	WI	68	415	no	
no					
115	ME	36	510	yes	
yes					
...
.					
3247	OK	146	510	no	
no					
3275	NY	120	510	no	
yes					
3290	CA	127	510	no	
no					
3291	MI	119	510	yes	
yes					
3310	NY	94	415	no	
no					

	Number Of Voicemail Messages	Total Day Minutes	Total Day Calls
\			
22	0	183.0	112
32	0	212.0	121
41	41	173.1	85
58	0	148.8	70
115	42	196.8	89

...
3247	0	138.4	104
3275	27	128.5	115
3290	0	107.9	128
3291	22	172.1	119
3310	0	190.4	91
	Total Day Charge	Total Evening Minutes	Total Evening Calls \
22	31.11	72.9	99
32	36.04	31.2	115
41	29.43	203.9	107
58	25.30	246.5	164
115	33.46	254.9	122
...
3247	23.53	158.9	122
3275	21.85	163.7	91
3290	18.34	187.0	77
3291	29.26	223.6	133
3310	32.37	92.0	107
	Total Evening Charge	Total Night Minutes	Total Night Calls \
22	6.20	181.8	78
32	2.65	293.3	78
41	17.33	122.2	78
58	20.95	129.8	103
115	21.67	138.3	126
...
3247	13.51	47.4	73
3275	13.91	242.9	121
3290	15.90	218.5	95
3291	19.01	150.0	94
3310	7.82	224.8	108
	Total Night Charge	Total International Minutes \	
22	8.18	9.5	
32	13.20	12.6	
41	5.50	14.6	
58	5.84	12.1	
115	6.22	20.0	
...	
3247	2.13	3.9	
3275	10.93	0.0	
3290	9.83	0.0	

3291	6.75	13.9
3310	10.12	13.6

	Total International Calls	Total International Charge \
22	19	2.57
32	10	3.40
41	15	3.94
58	3	3.27
115	6	5.40
...
3247	9	1.05
3275	0	0.00
3290	0	0.00
3291	20	3.75
3310	17	3.67

	Customer Service Calls	Churn
22	0	0
32	3	0
41	0	1
58	3	0
115	0	1
...
3247	4	1
3275	1	0
3290	0	0
3291	1	1
3310	2	0

[164 rows x 20 columns]

#Removing the outliers.

df1 = df[~outliers]

print("DataFrame after removing outliers:")

print(df1)

DataFrame after removing outliers:

Plan	State	Account Length	Area Code	International	Plan	Voice Mail
0	KS	128	415			no
yes						
1	OH	107	415			no
yes						
2	NJ	137	415			no
no						
3	OH	84	408			yes
no						
4	OK	75	415			yes
no						

...
3328	AZ	192	415	no	
yes					
3329	WV	68	415	no	
no					
3330	RI	28	510	no	
no					
3331	CT	184	510	yes	
no					
3332	TN	74	415	no	
yes					
	Number Of Voicemail Messages	Total Day Minutes	Total Day Calls		
\					
0		25	265.1	110	
1		26	161.6	123	
2		0	243.4	114	
3		0	299.4	71	
4		0	166.7	113	
...		
3328		36	156.2	77	
3329		0	231.1	57	
3330		0	180.8	109	
3331		0	213.8	105	
3332		25	234.4	113	
	Total Day Charge	Total Evening Minutes	Total Evening Calls	\	
0	45.07	197.4	99		
1	27.47	195.5	103		
2	41.38	121.2	110		
3	50.90	61.9	88		
4	28.34	148.3	122		
...		
3328	26.55	215.5	126		
3329	39.29	153.4	55		
3330	30.74	288.8	58		
3331	36.35	159.6	84		
3332	39.85	265.9	82		

	Total Evening Charge	Total Night Minutes	Total Night Calls \
0	16.78	244.7	91
1	16.62	254.4	103
2	10.30	162.6	104
3	5.26	196.9	89
4	12.61	186.9	121
...
3328	18.32	279.1	83
3329	13.04	191.3	123
3330	24.55	191.9	91
3331	13.57	139.2	137
3332	22.60	241.4	77

	Total Night Charge	Total International Minutes \
0	11.01	10.0
1	11.45	13.7
2	7.32	12.2
3	8.86	6.6
4	8.41	10.1
...
3328	12.56	9.9
3329	8.61	9.6
3330	8.64	14.1
3331	6.26	5.0
3332	10.86	13.7

	Total International Calls	Total International Charge \
0	3	2.70
1	3	3.70
2	5	3.29
3	7	1.78
4	3	2.73
...
3328	6	2.67
3329	4	2.59
3330	6	3.81
3331	10	1.35
3332	4	3.70

	Customer Service Calls	Churn
0	1	0
1	1	0
2	0	0
3	2	0
4	3	0
...
3328	2	0
3329	3	0
3330	2	0
3331	2	0

```
3332          0      0
```

```
[3169 rows x 20 columns]
```

```
df1.shape
```

```
(3169, 20)
```

After handling the outliers, we remain with 3169 rows, 164 rows removed.

EXPLORATORY DATA ANALYSIS

Let us see the statistical overview of the data.

```
df1.describe()
```

	Account Length	Number Of Voicemail Messages	Total Day Minutes
\count	3169.000000	3169.000000	3169.000000
mean	100.856737	7.974440	180.280120
std	39.474815	13.586481	53.835965
min	1.000000	0.000000	17.600000
25%	74.000000	0.000000	144.000000
50%	100.000000	0.000000	179.900000
75%	127.000000	19.000000	217.000000
max	217.000000	49.000000	337.400000

	Total Day Calls	Total Day Charge	Total Evening Minutes	\
count	3169.000000	3169.000000	3169.000000	
mean	100.60650	30.648157	201.086904	
std	19.72475	9.152086	50.080338	
min	42.00000	2.990000	49.200000	
25%	87.00000	24.480000	166.800000	
50%	101.00000	30.580000	201.400000	
75%	114.00000	36.890000	235.100000	
max	160.00000	57.360000	351.600000	

	Total Evening Calls	Total Evening Charge	Total Night Minutes
\count	3169.000000	3169.000000	3169.000000

mean	100.195014	17.092603	200.924361
std	19.614413	4.256827	49.990627
min	42.000000	4.180000	50.100000
25%	87.000000	14.180000	167.000000
50%	100.000000	17.120000	201.100000
75%	114.000000	19.980000	235.600000
max	159.000000	29.890000	352.500000

Total Night Calls			
Minutes \			
count	3169.000000	3169.000000	
mean	100.023982	9.041682	
std	19.405317	2.249643	
min	42.000000	2.250000	
25%	87.000000	7.520000	
50%	100.000000	9.050000	
75%	113.000000	10.600000	
max	158.000000	15.860000	

Total International Calls			
count	3169.000000	3169.000000	
mean	4.367939	2.784033	
std	2.156224	0.719748	
min	1.000000	0.540000	
25%	3.000000	2.300000	
50%	4.000000	2.780000	
75%	6.000000	3.270000	
max	11.000000	4.970000	

Customer Service Calls		Churn
count	3169.000000	3169.000000
mean	1.513411	0.139476
std	1.215649	0.346497
min	0.000000	0.000000
25%	1.000000	0.000000

50%	1.000000	0.000000
75%	2.000000	0.000000
max	5.000000	1.000000

The rows provide descriptive statistics including count, mean, standard deviation, minimum, 25th percentile, median, 75th percentile, and maximum values for each column in the dataset. We will then perform various analysis i.e univariate, bivariate and multivariate.

Univariate Analysis.

We will first look at the distribution of our target variable which Churn.


```
y = df1["Churn"].value_counts()
churn_percentage = y / len(df1) * 100
print(y)
print("")
print(churn_percentage)
sns.barplot(y.index, y.values)
```

```
0    2727
1     442
Name: Churn, dtype: int64
```

```
0    86.052382
1    13.947618
Name: Churn, dtype: float64
```

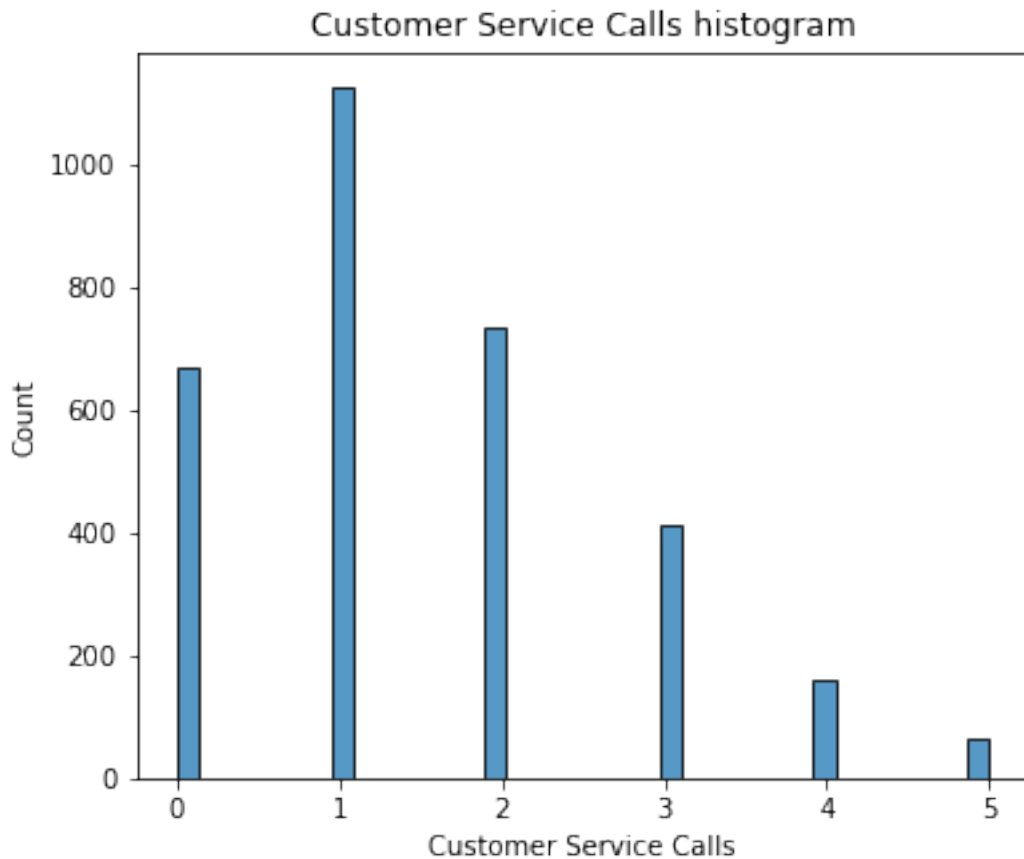
```
c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\
seaborn\_decorators.py:36: FutureWarning: Pass the following variables
as keyword args: x, y. From version 0.12, the only valid positional
argument will be `data`, and passing other arguments without an
explicit keyword will result in an error or misinterpretation.
  warnings.warn(
```

```
<AxesSubplot:>
```



This shows approximately 86% of the customers have not terminated the contract. We have a class imbalance for the target variable which could lead to biased predictions towards the no churn category. We will investigate into the use of oversampling when building the models.

```
#Distribution of the customer service calls.  
plt.figure(figsize=(6, 5))  
  
# Price boxplot  
  
sns.histplot(x=df1['Customer Service Calls'])  
plt.title('Customer Service Calls histogram')  
  
plt.show()
```



```
df1["Customer Service Calls"].value_counts()
```


```
1    1127
2     734
0     670
3     414
4     161
5      63
```

```
Name: Customer Service Calls, dtype: int64
```

Over 1000 customers have made 1 customer service call to the company. This is quite a huge number and might signify that there is one common problem. 0 customer service calls indicate that 670 customers are satisfied with the company's services. Although it is still a lower value compared to the customers who called once or twice.


#showing the distribution of the international column.

```
sns.countplot(x='International Plan', data=df1)
plt.title('Frequency of International Plan column')
plt.xlabel('Categories')
plt.ylabel('Frequency')
plt.show()
```




Less than 500 people have the international plan while almost 3000 people do not have the international plan.

```
#Visualizing the distribution of the Voice Mail plan.  
sns.countplot(x='Voice Mail Plan', data=df1)  
plt.title('Frequency of Voice Mail Plan')  
plt.xlabel('Categories')  
plt.ylabel('Frequency')  
plt.show()
```



More than 2000 customers do not have the Voice Mail Plan while less than 1000 people have the voice mail plan.

```
plt.figure(figsize=(10, 6))
sns.histplot(df1['Account Length'], bins=30, kde=True,
color='darkblue')
plt.title('Distribution of Account Length')
plt.xlabel('Account Length')
plt.ylabel('Frequency')
plt.show()
```




It shows that its distribution is a Gaussian distribution meaning it is symmetrical around the mean showing that data near the mean are more frequent in occurrence than data far from the mean.

Bivariate analysis.

We will visualize the relationship between Churn and State to see which states have the most and least churn rates.

```
df1.groupby(["State", "Churn"]).size().unstack().plot(kind='bar',
stacked=True, figsize=(30,10))
```

```
<AxesSubplot:xlabel='State'>
```



The above graph shows there are states with less rates of churn like AK, IA, LA and HI while there are states with higher rates of churn like TX, SC and NY.

Relationship between Churn and Account Length.

```
#calculating the median position for each churn category.
medians = df1.groupby('Churn')['Account Length'].median()
print("Medians:\n", medians)
```

```
Medians:
```

```
Churn
0      100.0
1      102.5
```

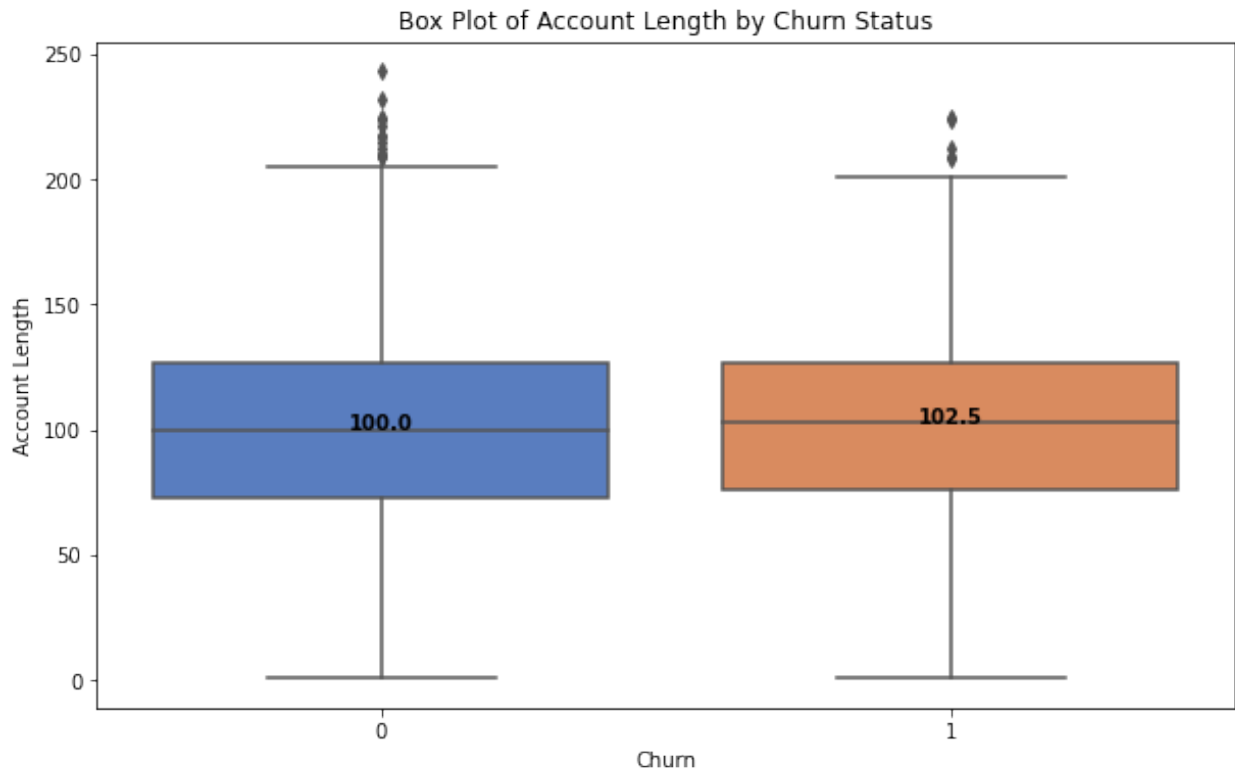
```
Name: Account Length, dtype: float64
```

```
plt.figure(figsize=(10, 6))
ax = sns.boxplot(x='Churn', y='Account Length', data=df,
palette='muted')
```

```
# Annotate the medians on the box plot
```

```
for tick, label in enumerate(ax.get_xticklabels()):
    x_position = tick
    y_position = medians[tick]
    ax.text(x_position, y_position, f'{y_position:.1f}',
            horizontalalignment='center', size='medium',
color='black', weight='semibold')
```

```
plt.title('Box Plot of Account Length by Churn Status')
plt.xlabel('Churn')
plt.ylabel('Account Length')
plt.show()
```

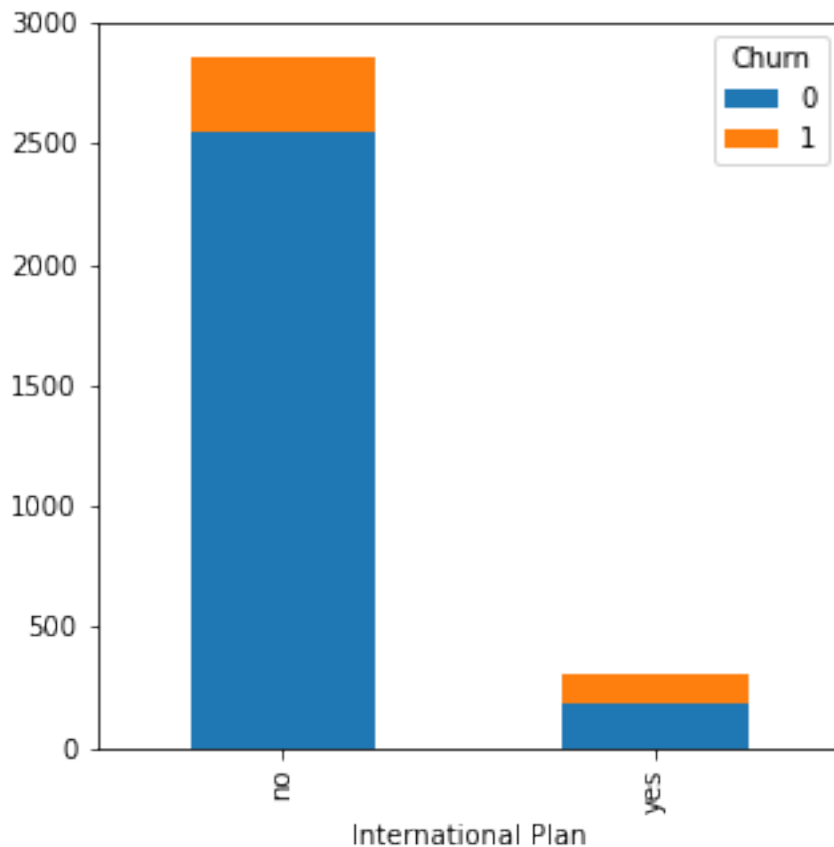



There is a slight difference in the medians of the churn categories. The customers who churn have a median of 102.5, which shows that the customers with longer account lengths stop using the company's services. But when we look at the outliers, the category false has more outliers which shows a higher variability in the account length, showing that the values are less consistent.

Relationship between Churn and International plan.

```
df1.groupby(["International Plan",  
"Churn"]).size().unstack().plot(kind='bar', stacked=True,  
figsize=(5,5))
```

```
<AxesSubplot:xlabel='International Plan'>
```

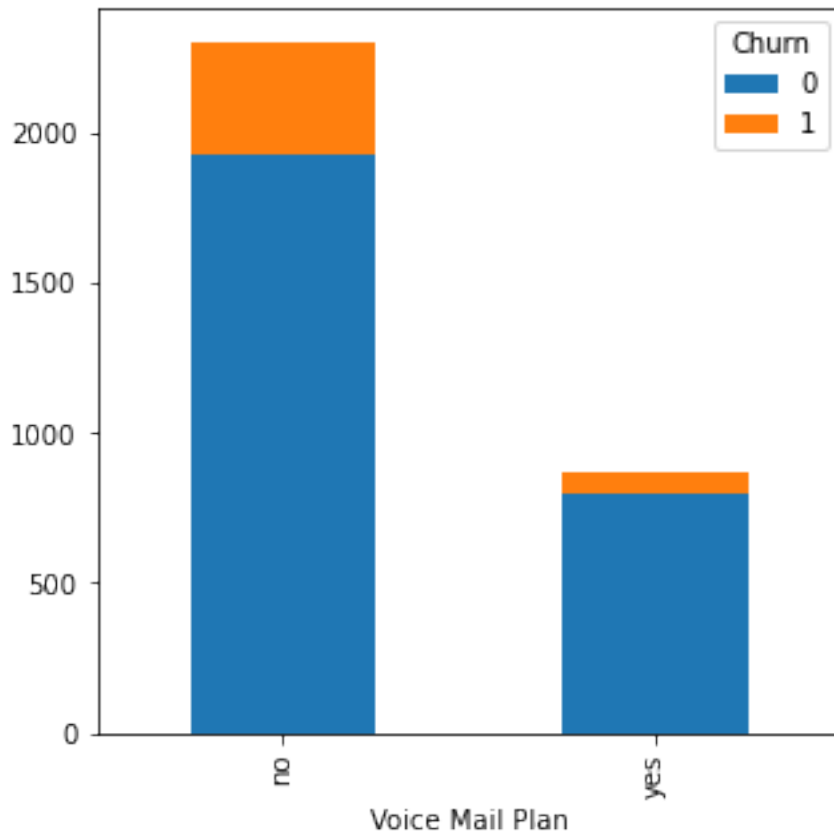


Both with or without the international plan have churn rates. Interestingly, those without the international plan have a much higher rate of churn compared to those with an international plan.

Relationship between churn and voice mail plan

```
df1.groupby(["Voice Mail Plan",  
"Churn"]).size().unstack().plot(kind='bar', stacked=True,  
figsize=(5,5))
```

```
<AxesSubplot:xlabel='Voice Mail Plan'>
```

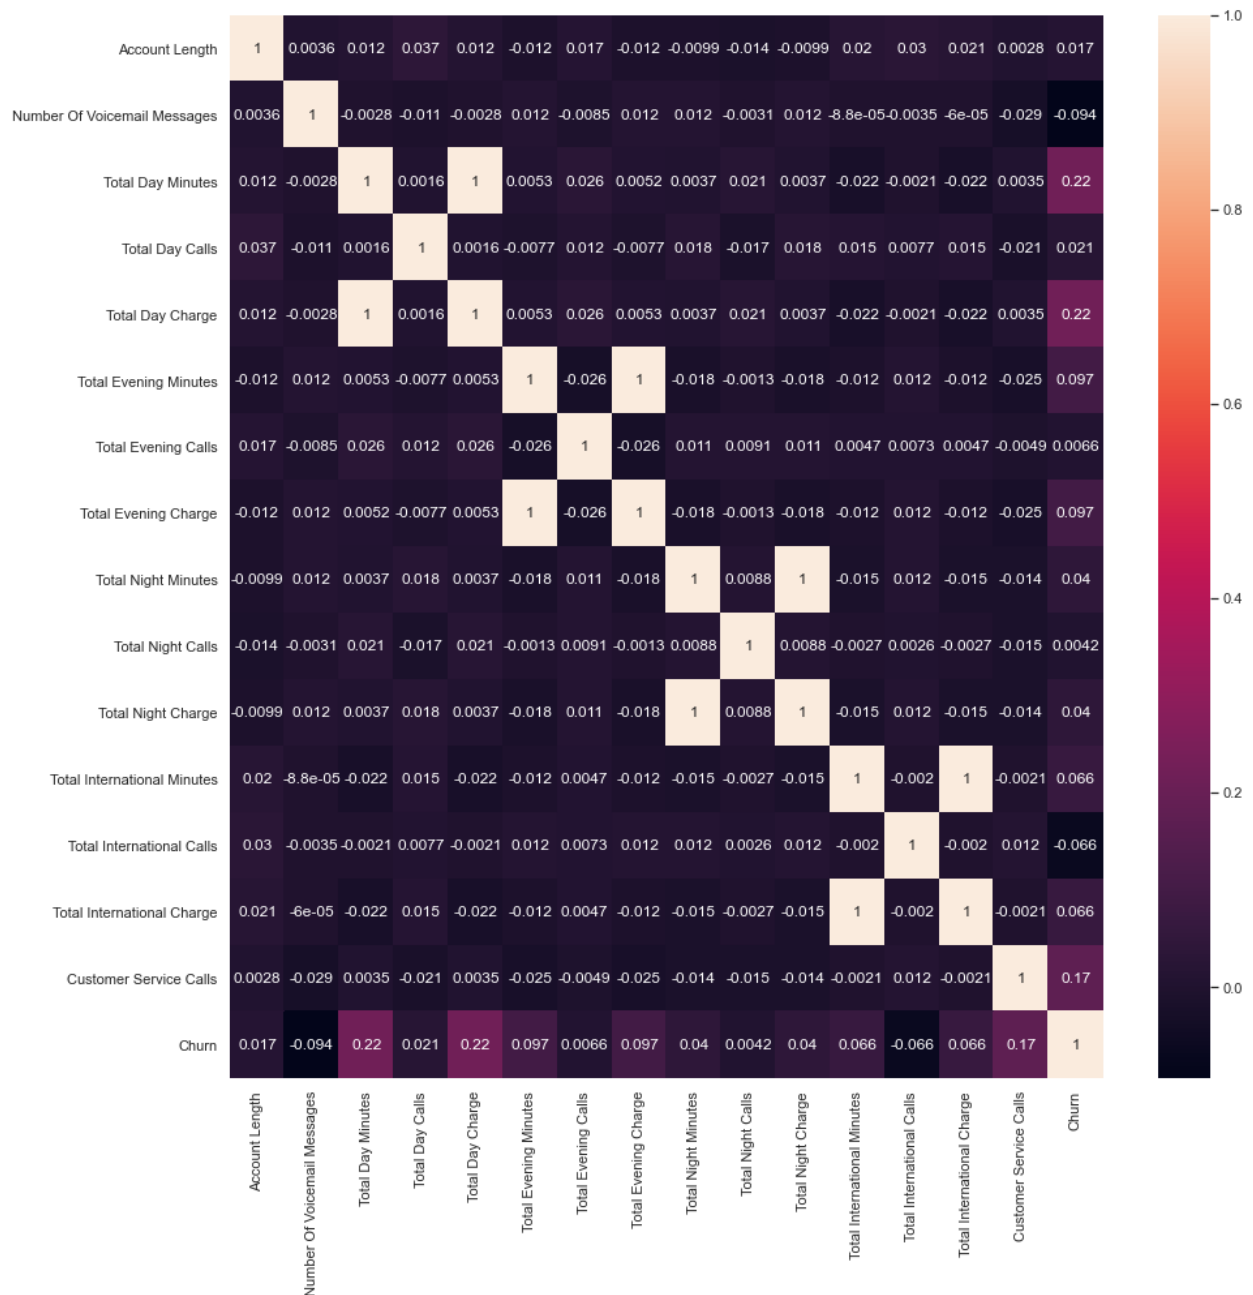


Almost similar to the previous graph, the customers without a voice mail plan have a higher rate of churn compared to those with a voicemail plan.

Multivariate analysis.

Creating a heat map to show correlation of the entire dataset, in which high correlations are coloured more to the red and lower ones more to purple.

```
sns.set(rc={'figure.figsize':(15, 15)})  
#Use the .heatmap method to depict the relationship visually  
sns.heatmap(df1.corr(), annot=True, annot_kws={"size": 12})  
plt.show()
```



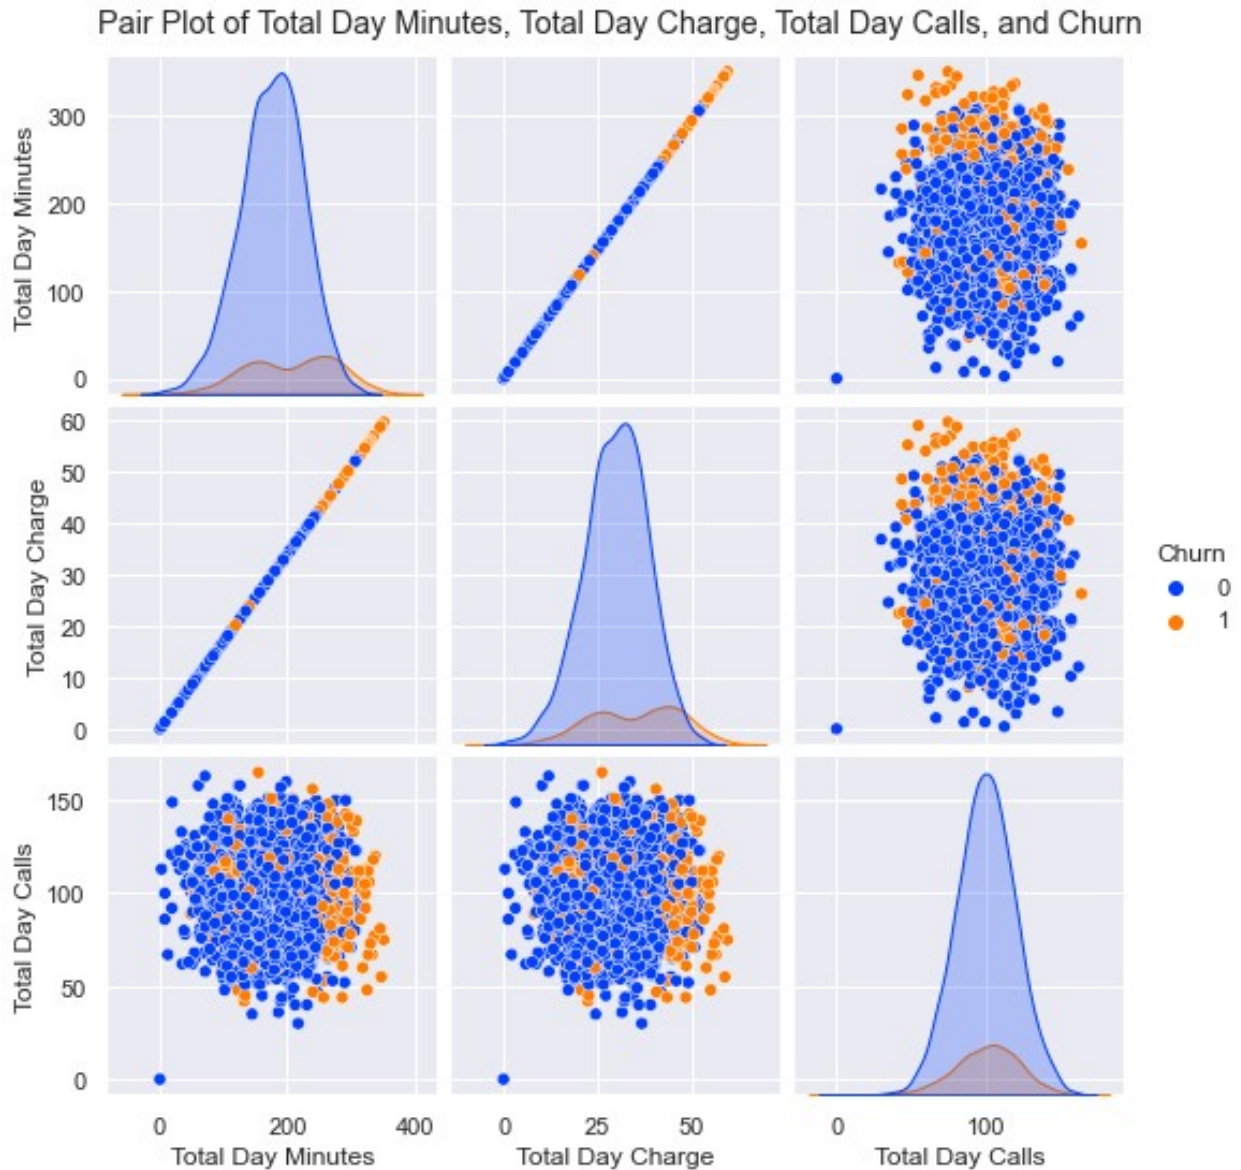
From the above , we can see the columns that highly correlated to churn are Customer Service Calls, total day minutes and total day charge.

Let us now visualize the correlation between different columns.

```
#visualizing the relationship between total day minutes, total day
charge , total day calls and churn
plt.figure(figsize=(12, 10))
sns.pairplot(df, vars=['Total Day Minutes', 'Total Day Charge', 'Total
Day Calls'], hue='Churn', palette='bright')
plt.suptitle('Pair Plot of Total Day Minutes, Total Day Charge, Total
```

```
Day Calls, and Churn', y=1.02)
plt.show()
```

<Figure size 864x720 with 0 Axes>

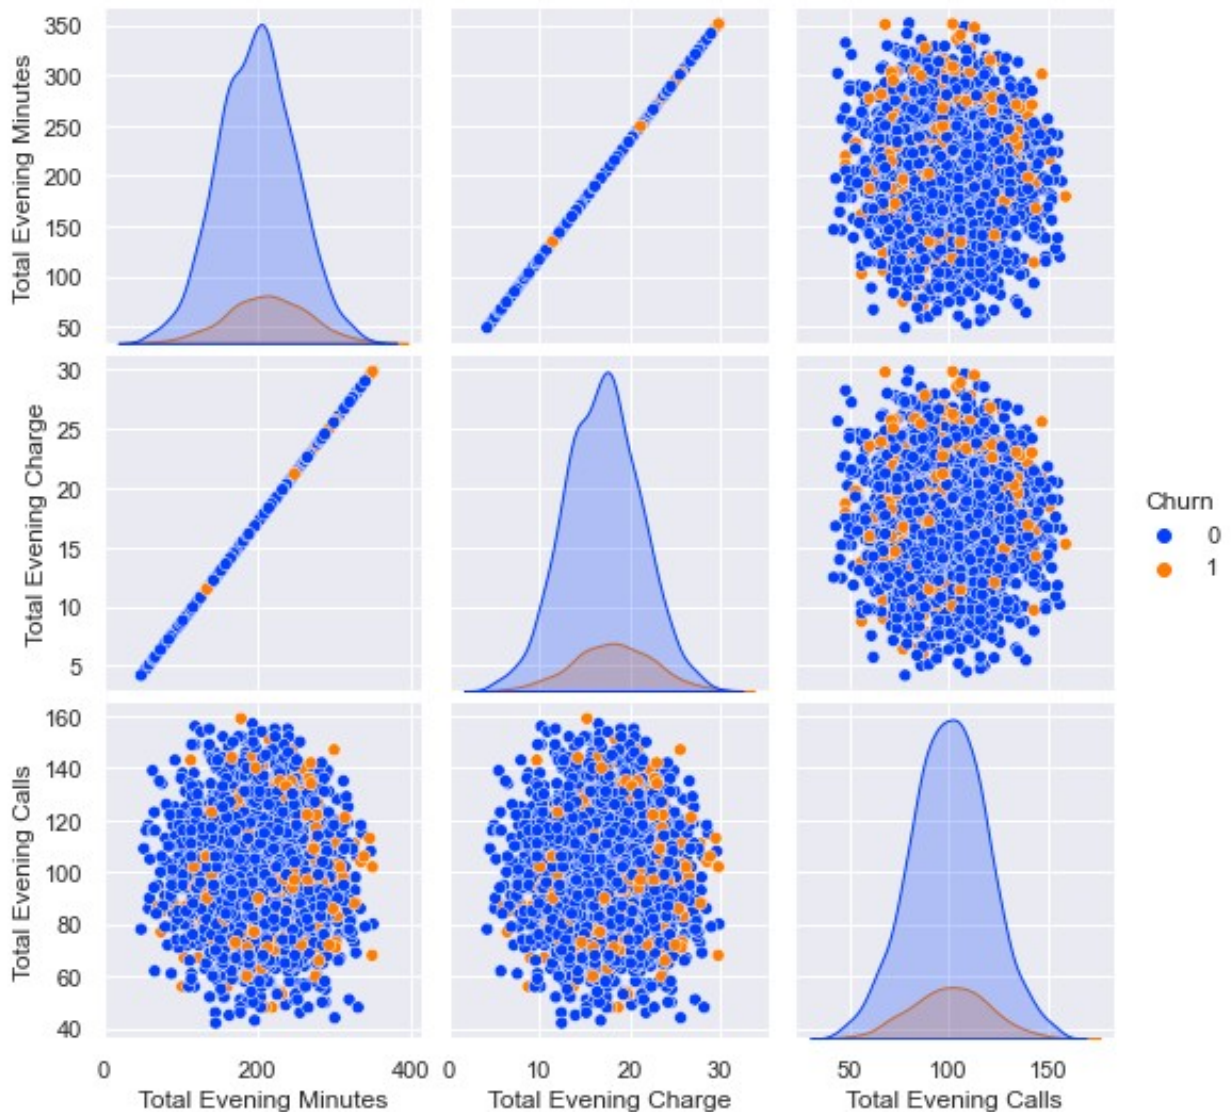


We see that there is a linear relationship between total day charge and total day minutes charge, in relation with the churn rate the higher day charge show more customers tend to leave the company. The cluster points show higher churn rate when the total day minutes and total day charge increase.

```
#visualizing the relationship between total evening minutes, total
evening charge, total evening calls and churn
sns.pairplot(df1, vars=['Total Evening Minutes', 'Total Evening
Charge', 'Total Evening Calls'], hue='Churn', palette='bright',
```

```
diag_kind='kde')
plt.suptitle('Pair Plot of Total Evening Minutes, Total Evening
Charges, Total Evening Calls and Churn', y=1.02)
plt.show()
```

Pair Plot of Total Evening Minutes, Total Evening Charges, Total Evening Calls and Churn



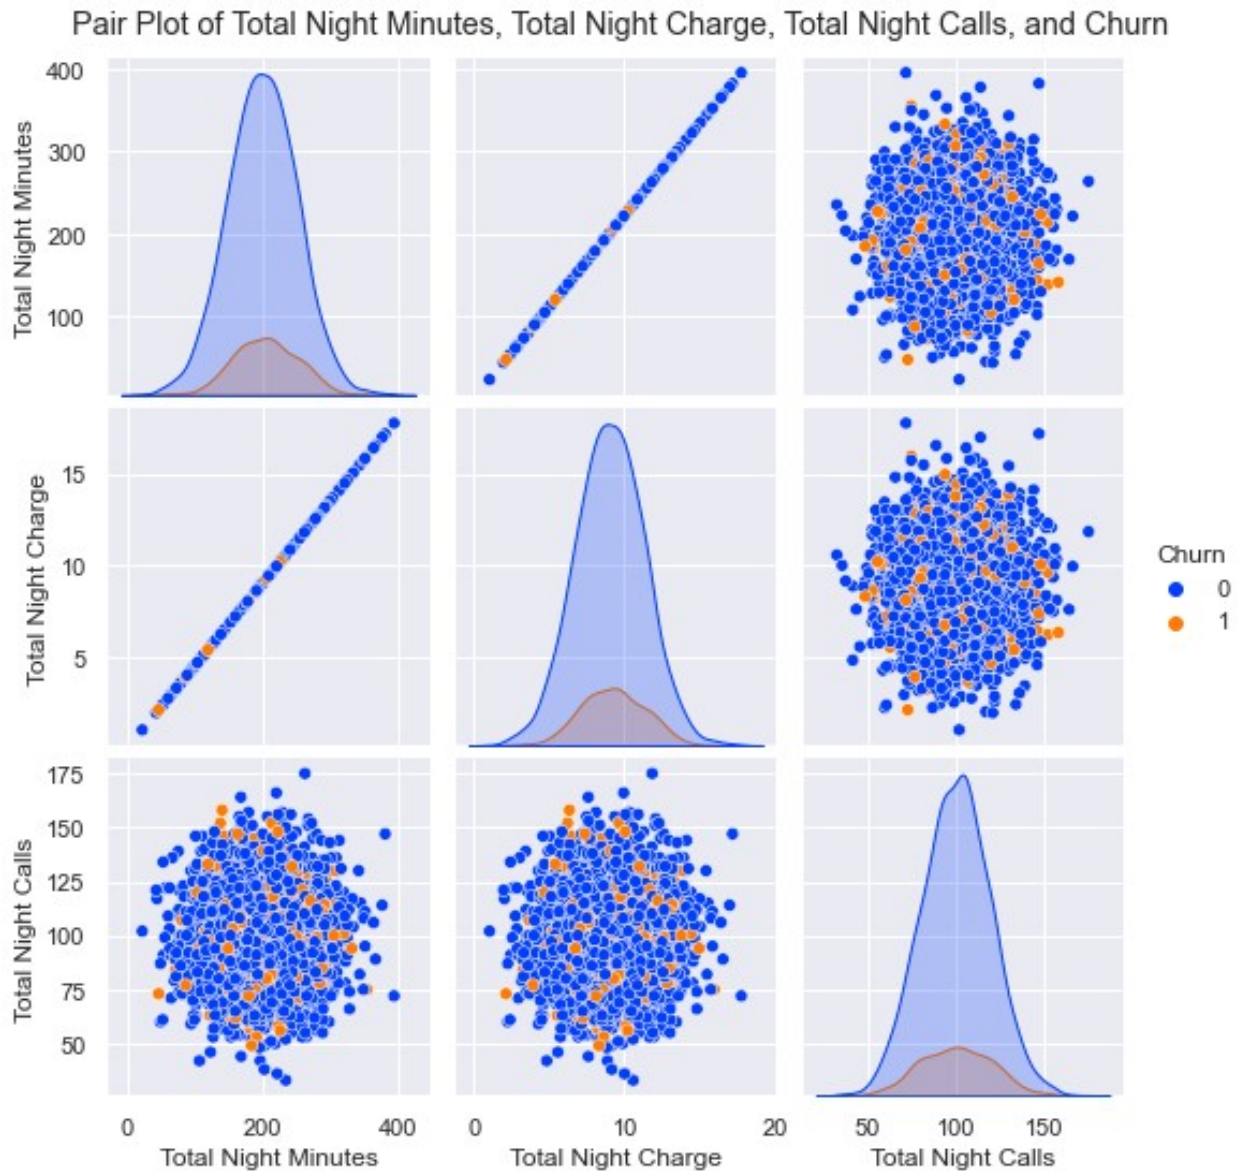
There is a clear linear relationship between the total evening charges and total evening minutes. The cluster points show a low churn rate and seem to be unevenly distributed. It spread across the different columns.

```
#visualizing the relationship between total night calls, total night
charge, total night minutes and churn
plt.figure(figsize=(12, 10))
sns.pairplot(df, vars=['Total Night Minutes', 'Total Night Charge',
'Total Night Calls'], hue='Churn', palette='bright')
```



```
plt.suptitle('Pair Plot of Total Night Minutes, Total Night Charge,  
Total Night Calls, and Churn', y=1.02)  
plt.show()
```

<Figure size 864x720 with 0 Axes>



There is a clear linear relationship between the total night charges and total night minutes. The cluster points show a low churn rate and seem to be unevenly distributed. It spread across the different columns.

Feature Engineering.

We will create new columns from different columns.

#We will create new columns namely, total calls, average call duration, customer tenure and total charges. Customer #tenure refers to the average time a customer stays in business with a company. It will be calculated in months.

```
df1.loc[:, 'Total Charges'] = df1['Total Day Charge'] + df1['Total Evening Charge'] + df1['Total Night Charge'] + df1['Total International Charge']
df1.loc[:, 'Total Calls'] = df1['Total Day Calls'] + df1['Total Evening Calls'] + df1['Total Night Calls'] + df1['Total International Calls']
df1.loc[:, 'Average Call Duration'] = (df1['Total Day Minutes'] + df1['Total Evening Minutes'] + df1['Total Night Minutes'] + df1['Total International Minutes'])/4
df1.loc[:, 'Customer Tenure'] = df1['Account Length'] / 30
```

```
c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\pandas\core\indexing.py:1596: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
self.obj[key] = _infer_fill_value(value)
```

```
c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\pandas\core\indexing.py:1745: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead
```

See the caveats in the documentation:

https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
isetter(ilocs[0], value)
```

```
df1.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
Int64Index: 3169 entries, 0 to 3332
```

```
Data columns (total 24 columns):
```

#	Column	Non-Null Count	Dtype
0	State	3169 non-null	object
1	Account Length	3169 non-null	int64
2	Area Code	3169 non-null	object
3	International Plan	3169 non-null	object
4	Voice Mail Plan	3169 non-null	object
5	Number Of Voicemail Messages	3169 non-null	int64
6	Total Day Minutes	3169 non-null	float64
7	Total Day Calls	3169 non-null	int64
8	Total Day Charge	3169 non-null	float64

9	Total Evening Minutes	3169	non-null	float64
10	Total Evening Calls	3169	non-null	int64
11	Total Evening Charge	3169	non-null	float64
12	Total Night Minutes	3169	non-null	float64
13	Total Night Calls	3169	non-null	int64
14	Total Night Charge	3169	non-null	float64
15	Total International Minutes	3169	non-null	float64
16	Total International Calls	3169	non-null	int64
17	Total International Charge	3169	non-null	float64
18	Customer Service Calls	3169	non-null	int64
19	Churn	3169	non-null	int32
20	Total Charges	3169	non-null	float64
21	Total Calls	3169	non-null	int64
22	Average Call Duration	3169	non-null	float64
23	Customer Tenure	3169	non-null	float64

dtypes: float64(11), int32(1), int64(8), object(4)
memory usage: 766.6+ KB

Pre-processing data.

This includes splitting the dataset, normalize/standardize the data, performing one hot encoding and label encoding and addressing multicollinearity.

```
from sklearn.model_selection import train_test_split

X = df1.drop('Churn', axis=1)
y = df1['Churn']
X_train, X_test, y_train, y_test = train_test_split(X, y,
test_size=0.2, random_state=42)

from sklearn.preprocessing import LabelEncoder, OneHotEncoder
from sklearn.compose import ColumnTransformer

# Label encode the State and Area Code columns
le = LabelEncoder()
X_train['Encoded State'] = le.fit_transform(X_train['State'])
X_test['Encoded State'] = le.transform(X_test['State'])
X_train['Encoded Area Code'] = le.fit_transform(X_train['Area Code'])
X_test['Encoded Area Code'] = le.transform(X_test['Area Code'])

# One-hot encode the voicemail plan and international plan columns
ohe = OneHotEncoder(drop='first', sparse=False)

# Fit and transform the voice mail plan and international_plan columns
on the training set
onehot_encoded_train = ohe.fit_transform(X_train[['Voice Mail Plan',
'International Plan']])
onehot_encoded_test = ohe.transform(X_test[['Voice Mail Plan',
```

```

'International Plan']])

# Convert the one-hot encoded columns back to DataFrame
onehot_encoded_train_df = pd.DataFrame(onehot_encoded_train,
columns=ohe.get_feature_names_out(['Voice Mail Plan', 'International
Plan']))
onehot_encoded_test_df = pd.DataFrame(onehot_encoded_test,
columns=ohe.get_feature_names_out(['Voice Mail Plan', 'International
Plan']))

# Reset index to match original DataFrame after split
onehot_encoded_train_df.reset_index(drop=True, inplace=True)
onehot_encoded_test_df.reset_index(drop=True, inplace=True)
X_train.reset_index(drop=True, inplace=True)
X_test.reset_index(drop=True, inplace=True)

# Drop the original state, Area code voice mail plan, and
international plan columns
X_train.drop(['State', 'Area Code', 'Voice Mail Plan', 'International
Plan'], axis=1, inplace=True)
X_test.drop(['State', 'Area Code', 'Voice Mail Plan', 'International
Plan'], axis=1, inplace=True)

# Concatenate the new one-hot encoded columns with the original
DataFrame
X_train_final = pd.concat([X_train, onehot_encoded_train_df], axis=1)
X_test_final = pd.concat([X_test, onehot_encoded_test_df], axis=1)

```

```

c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\
sklearn\preprocessing\_encoders.py:975: FutureWarning: `sparse` was
renamed to `sparse_output` in version 1.2 and will be removed in 1.4.
`sparse_output` is ignored unless you leave `sparse` to its default
value.

```

```
warnings.warn(
```

```
X_train_final.head()
```

	Account Length	Number Of Voicemail Messages	Total Day Minutes	\
0	45	0	159.8	
1	119	19	178.1	
2	104	0	183.6	
3	125	0	298.4	
4	84	0	216.1	

	Total Day Calls	Total Day Charge	Total Evening Minutes	\
0	91	27.17	120.4	
1	110	30.28	212.8	
2	133	31.21	120.7	

3	78	50.73	270.5
4	114	36.74	197.5
Total Evening Calls Total Evening Charge Total Night Minutes \			
0	86	10.23	163.0
1	100	18.09	226.3
2	98	10.26	215.1
3	142	22.99	107.3
4	107	16.79	217.8
Total Night Calls ... Total International Charge Customer Service Calls \			
0	93	...	2.86
2			
1	123	...	2.70
1			
2	112	...	3.43
1			
3	84	...	3.29
0			
4	104	...	2.65
1			
Total Charges Total Calls Average Call Duration Customer Tenure \			
0	47.60	273	113.450 1.500000
1	61.25	339	156.800 3.966667
2	54.58	345	133.025 3.466667
3	81.84	306	172.100 4.166667
4	65.98	328	160.300 2.800000
Encoded State Encoded Area Code Voice Mail Plan_yes \			
0	12	2	0.0
1	30	1	1.0
2	50	0	0.0
3	30	1	0.0
4	23	1	0.0
International Plan_yes			
0		0.0	
1		0.0	
2		0.0	
3		1.0	
4		0.0	

[5 rows x 23 columns]

X_test_final.head()

	Account Length	Number Of Voicemail Messages	Total Day Minutes	\
0	155	0	262.4	
1	96	0	106.6	
2	87	0	146.3	
3	133	0	277.3	
4	162	0	220.6	

	Total Day Calls	Total Day Charge	Total Evening Minutes	\
0	55	44.61	194.6	
1	128	18.12	284.8	
2	108	24.87	171.8	
3	138	47.14	228.4	
4	117	37.50	155.2	

	Total Evening Calls	Total Evening Charge	Total Night Minutes	\
0	113	16.54	146.5	
1	87	24.21	178.9	
2	102	14.60	167.5	
3	117	19.41	117.3	
4	121	13.19	186.7	

	Total Night Calls	...	Total International Charge	Customer Service Calls	\
0	85	...	2.24		
2					
1	92	...	4.02		
1					
2	66	...	1.43		
1					
3	103	...	3.46		
2					
4	89	...	2.84		
1					

	Total Charges	Total Calls	Average Call Duration	Customer Tenure	\
0	69.98	259	152.950	5.166667	
1	54.40	314	146.300	3.200000	
2	48.44	285	122.725	2.900000	
3	75.29	362	158.950	4.433333	
4	61.93	338	143.250	5.400000	

	Encoded State	Encoded Area Code	Voice Mail Plan_yes \
0	27	0	0.0
1	35	1	0.0
2	44	1	0.0
3	22	0	0.0
4	46	2	0.0

	International Plan_yes
0	0.0
1	0.0
2	0.0
3	0.0
4	0.0

[5 rows x 23 columns]

Standardizing the data.

From the describe function, we can see that our data is not on a consistent scale of 0 to 1. We will use the StandardScaler to normalize the data. Ensures a mean of 0 and standard deviation of 1.

```
from sklearn.preprocessing import StandardScaler

#intialize the scaler
scaler = StandardScaler()

# Fit and transform the training set
X_train_scaled = scaler.fit_transform(X_train_final)

# Transform the test set
X_test_scaled = scaler.transform(X_test_final)

# Convert the scaled data back to DataFrame
X_train_df = pd.DataFrame(X_train_scaled,
                           columns=X_train_final.columns)
X_test_df = pd.DataFrame(X_test_scaled, columns=X_test_final.columns)

X_train_df.describe()
```

	Account Length	Number Of Voicemail Messages	Total Day Minutes
\			
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	-1.709787e-16	4.204395e-17	1.541611e-17
std	1.000197e+00	1.000197e+00	1.000197e+00
min	-2.531471e+00	-5.874908e-01	-3.022368e+00

25%	-6.786879e-01	-5.874908e-01	-6.686012e-01
50%	6.587939e-03	-5.874908e-01	-8.432386e-03
75%	6.664832e-01	8.083504e-01	6.758776e-01
max	2.950736e+00	2.938845e+00	2.916366e+00
	Total Day Calls	Total Day Charge	Total Evening Minutes \
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	-2.985120e-16	-2.074168e-16	8.128497e-17
std	1.000197e+00	1.000197e+00	1.000197e+00
min	-2.962570e+00	-3.022653e+00	-3.023636e+00
25%	-6.842309e-01	-6.686067e-01	-6.841885e-01
50%	2.458590e-02	-8.818209e-03	1.824114e-02
75%	6.827729e-01	6.760947e-01	6.690799e-01
max	3.011742e+00	2.916535e+00	2.976780e+00
	Total Evening Calls	Total Evening Charge	Total Night Minutes
\			
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	-2.858989e-16	-1.653729e-16	-2.102197e-16
std	1.000197e+00	1.000197e+00	1.000197e+00
min	-2.936851e+00	-3.024156e+00	-3.023809e+00
25%	-6.572674e-01	-6.850550e-01	-6.786380e-01
50%	1.278925e-03	1.760908e-02	4.743154e-03
75%	6.598252e-01	6.689156e-01	6.921383e-01
max	2.990066e+00	2.977669e+00	3.045338e+00
	Total Night Calls	...	Total International Charge \
count	2.535000e+03	...	2.535000e+03
mean	-1.023069e-16	...	-1.289348e-16
std	1.000197e+00	...	1.000197e+00
min	-3.008707e+00	...	-3.126444e+00
25%	-6.700071e-01	...	-6.541066e-01
50%	5.617391e-03	...	2.648052e-02
75%	6.812419e-01	...	6.653990e-01
max	3.019942e+00	...	3.026620e+00
	Customer Service Calls	Total Charges	Total Calls \
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	-7.567911e-17	-5.269508e-16	2.662783e-17

std	1.000197e+00	1.000197e+00	1.000197e+00
min	-1.244978e+00	-3.286807e+00	-3.330667e+00
25%	-4.200907e-01	-6.795813e-01	-6.809919e-01
50%	-4.200907e-01	-8.738943e-03	7.045241e-03
75%	4.047970e-01	6.664408e-01	6.804432e-01
max	2.879460e+00	3.148847e+00	3.022697e+00

	Average Call Duration	Customer Tenure	Encoded State \
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	4.092278e-16	3.573736e-16	-5.746006e-17
std	1.000197e+00	1.000197e+00	1.000197e+00
min	-3.477911e+00	-2.531471e+00	-1.773496e+00
25%	-6.622008e-01	-6.786879e-01	-8.218433e-01
50%	1.652125e-02	6.587939e-03	-6.140563e-03
75%	6.687703e-01	6.664832e-01	8.775374e-01
max	3.257489e+00	2.950736e+00	1.625265e+00

	Encoded Area Code	Voice Mail Plan_yes	International Plan_yes
count	2.535000e+03	2.535000e+03	2.535000e+03
mean	5.710970e-17	6.446739e-17	7.007325e-18
std	1.000197e+00	1.000197e+00	1.000197e+00
min	-1.432145e+00	-6.139780e-01	-3.263491e-01
25%	-2.443365e-02	-6.139780e-01	-3.263491e-01
50%	-2.443365e-02	-6.139780e-01	-3.263491e-01
75%	1.383278e+00	1.628723e+00	-3.263491e-01
max	1.383278e+00	1.628723e+00	3.064204e+00

[8 rows x 23 columns]

X_test_df.describe()

	Account Length	Number Of Voicemail Messages	Total Day Minutes
count	634.000000	634.000000	634.000000
mean	0.014755	-0.008227	-0.006865
std	1.009357	0.990605	0.998689
min	-2.480710	-0.587491	-2.615681
25%	-0.678688	-0.587491	-0.702956

50%	-0.018793	-0.587491	-0.007504
75%	0.685519	0.789984	0.695841
max	2.899975	3.012310	2.743664

	Total Day Calls	Total Day Charge	Total Evening Minutes \
count	634.000000	634.000000	634.000000
mean	0.023308	-0.006863	-0.048977
std	0.993054	0.998689	0.967202
min	-2.861311	-2.615201	-2.724013
25%	-0.633601	-0.702470	-0.731315
50%	0.075216	-0.007726	-0.071051
75%	0.682773	0.695757	0.638324
max	2.606704	2.743942	2.522383

	Total Evening Calls	Total Evening Charge	Total Night Minutes
\			
count	634.000000	634.000000	634.000000
mean	0.055772	-0.048966	0.016120
std	0.966336	0.967207	1.016347
min	-2.734221	-2.723014	-2.875292
25%	-0.657267	-0.731744	-0.673620
50%	0.102594	-0.071099	0.023810
75%	0.761140	0.637984	0.734787
max	2.736779	2.522455	3.039317

	Total Night Calls ...	Total International Charge \
count	634.000000 ...	634.000000
mean	0.034308 ...	-0.047918
std	1.041470 ...	0.997560
min	-3.008707 ...	-3.084775
25%	-0.670007 ...	-0.723554
50%	0.057589 ...	-0.084636
75%	0.733213 ...	0.623730
max	2.864029 ...	2.943282

	Customer Service Calls	Total Charges	Total Calls \
count	634.000000	634.000000	634.000000
mean	0.017074	-0.026108	0.063523
std	1.013692	0.991107	0.983035

min	-1.244978	-3.142228	-2.803659
25%	-0.420091	-0.727051	-0.600477
50%	-0.420091	-0.019823	0.065602
75%	0.404797	0.688850	0.680443
max	2.879460	2.939691	2.876306

	Average Call Duration	Customer Tenure	Encoded State \
count	634.000000	634.000000	634.000000
mean	-0.024364	0.014755	-0.028334
std	0.992871	1.009357	1.034717
min	-3.151224	-2.480710	-1.773496
25%	-0.725004	-0.678688	-0.957794
50%	-0.012205	-0.018793	-0.006141
75%	0.670742	0.685519	0.928519
max	2.764078	2.899975	1.625265

	Encoded Area Code	Voice Mail Plan_yes	International Plan_yes
count	634.000000	634.000000	634.000000
mean	-0.091045	-0.005548	0.021263
std	0.998668	0.997953	1.029288
min	-1.432145	-0.613978	-0.326349
25%	-1.432145	-0.613978	-0.326349
50%	-0.024434	-0.613978	-0.326349
75%	-0.024434	1.628723	-0.326349
max	1.383278	1.628723	3.064204

[8 rows x 23 columns]

Multicollinearity

In machine learning, multicollinearity is not such a big deal since it should not mess with the predictive power in classification models hence we will not handle multicollinearity.

Feature Selection

This process is important and is used to make the machine learning process accurate. It increases the prediction power by selecting the most critical values and eliminating the redundant ones. We will use a filter method called Variance threshold.

```
from sklearn.feature_selection import VarianceThreshold
# Specify the threshold value
```

```

threshold_value = 0.5

# Initialize the variance threshold selector with the specified
# threshold
selector = VarianceThreshold(threshold=threshold_value)

# Fit the selector on the training data
selector.fit(X_train_df)

# Get the indices of the features with non-zero variance
selected_indices = selector.get_support(indices=True)

# Get the selected features
selected_features = X_train_df.columns[selected_indices]

# Create a new DataFrame with the selected features
X_train_new= X_train_df[selected_features]
X_test_new = X_test_df[selected_features]

# Display the selected features
print("Selected features after variance thresholding:")
X_test_new.head()

```

Selected features after variance thresholding:

	Account Length	Number Of Voicemail Messages	Total Day Minutes	\
0	1.377140	-0.587491	1.523605	
1	-0.120315	-0.587491	-1.369624	
2	-0.348740	-0.587491	-0.632389	
3	0.818767	-0.587491	1.800300	
4	1.554804	-0.587491	0.747373	

	Total Day Calls	Total Day Charge	Total Evening Minutes	\
0	-2.304383	1.523770	-0.138516	
1	1.391590	-1.369905	1.651291	
2	0.378994	-0.632559	-0.590928	
3	1.897887	1.800138	0.532166	
4	0.834662	0.747098	-0.920316	

	Total Evening Calls	Total Evening Charge	Total Night Minutes	\
0	0.659825	-0.138798	-1.089068	
1	-0.657267	1.651711	-0.438802	
2	0.102594	-0.591678	-0.667599	
3	0.862455	0.531184	-1.675110	
4	1.065084	-0.920833	-0.282257	

	Total Night Calls	...	Total International Charge	Customer
Service Calls	\			
0	-0.773949	...	-0.765223	
0.404797				
1	-0.410152	...	1.707114	-

```

0.420091
2      -1.761401    ...      -1.890275      -
0.420091
3      0.161531    ...      0.929300
0.404797
4      -0.566065    ...      0.068149      -
0.420091

  Total Charges  Total Calls  Average Call Duration  Customer Tenure
\
0      0.998489    -1.339751      0.211407      1.377140
1      -0.503196    0.270549      -0.088244      -0.120315
2      -1.077653    -0.578518      -1.150543      -0.348740
3      1.510295     1.675901      0.481770      0.818767
4      0.222586     0.973225      -0.225678      1.554804

  Encoded State  Encoded Area Code  Voice Mail Plan_yes  \
0      0.061835      -1.432145      -0.613978
1      0.605636      -0.024434      -0.613978
2      1.217414      -0.024434      -0.613978
3      -0.278041     -1.432145      -0.613978
4      1.353364      1.383278      -0.613978

  International Plan_yes
0      -0.326349
1      -0.326349
2      -0.326349
3      -0.326349
4      -0.326349

[5 rows x 23 columns]

```

With the variance threshold of 0.5, it seems all features are relevant for modeling.

Modeling

We will start with a baseline model and go ahead and build more models and compare their performances.


Baseline metrics.

Before we perform any modeling, let's start with a dummy classifier that always predicts the positive class. This will be useful for detecting imbalanced classes by providing a comparison point. We will focus on the test data.

```
np.bincount(y_test)
array([536,  98], dtype=int64)
```

We will then calculate using different classification metrics to evaluate the model's performance from both positive and negative values.

```
#The confusion matrix
from sklearn.metrics import confusion_matrix, ConfusionMatrixDisplay
majority_class = y_train.mode()[0]
y_pred_baseline = [majority_class] * len(y_test)
conf_matrix = confusion_matrix(y_test, y_pred_baseline)
ConfusionMatrixDisplay(conf_matrix, display_labels=['Non-Churn',
'Churn']).plot(cmap='Greens')
plt.title('Confusion Matrix - Baseline Metrics')
plt.show()
```



```
#Calculate the baseline accuracy, precision, flscore and recall.
from sklearn.metrics import accuracy_score, precision_score,
recall_score, f1_score, classification_report
#evaluate the baseline model
baseline_accuracy = accuracy_score(y_test, y_pred_baseline)
baseline_precision = precision_score(y_test, y_pred_baseline,
pos_label=majority_class)
baseline_recall = recall_score(y_test, y_pred_baseline,
pos_label=majority_class)
baseline_f1 = f1_score(y_test, y_pred_baseline,
pos_label=majority_class)

report = classification_report(y_test, y_pred_baseline)
```

```

print(f'Baseline Model Metrics:')
print(f'Accuracy: {baseline_accuracy}')
print(f'Precision: {baseline_precision}')
print(f'Recall: {baseline_recall}')
print(f'F1-Score: {baseline_f1}')

```

```

print('Classification Report:')
print(report)

```

Baseline Model Metrics:

Accuracy: 0.8454258675078864

Precision: 0.8454258675078864

Recall: 1.0

F1-Score: 0.9162393162393162

Classification Report:

	precision	recall	f1-score	support
0	0.85	1.00	0.92	536
1	0.00	0.00	0.00	98
accuracy			0.85	634
macro avg	0.42	0.50	0.46	634
weighted avg	0.71	0.85	0.77	634

```

c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\
sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control
this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))

```

```

c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\
sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control
this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))

```

```

c:\Users\Administrator\anaconda3\envs\learn-env\lib\site-packages\
sklearn\metrics\_classification.py:1471: UndefinedMetricWarning:
Precision and F-score are ill-defined and being set to 0.0 in labels
with no predicted samples. Use `zero_division` parameter to control
this behavior.

```

```

_warn_prf(average, modifier, msg_start, len(result))

```

Logistic Regression

Instantiate and Fit the Logistic Regression model.

```

#importing the class
from sklearn.linear_model import LogisticRegression
#Instantiate the model
model = LogisticRegression(fit_intercept=False, C=1e12,
solver='liblinear', random_state = 42)
#fitting the model on scaled data
model.fit(X_train_new, y_train)
# evaluate performance on the train set
y_pred_train = model.predict(X_train_new)

train_residuals = np.abs(y_train - y_pred_train)
print(pd.Series(train_residuals, name="Residuals
(counts)").value_counts())
print()
print(pd.Series(train_residuals, name="Residuals
(proportions)").value_counts(normalize=True))

0    1649
1     886
Name: Residuals (counts), dtype: int64

0    0.650493
1    0.349507
Name: Residuals (proportions), dtype: float64

```

Our model was about 65% correct on the training data.

```

#evaluate performance on the test set.
y_pred_test = model.predict(X_test_new)

test_residuals = np.abs(y_test - y_pred_test)
print(pd.Series(test_residuals, name="Residuals
(counts)").value_counts())
print()
print(pd.Series(test_residuals, name="Residuals
(proportions)").value_counts(normalize=True))

0    419
1    215
Name: Residuals (counts), dtype: int64

0    0.660883
1    0.339117
Name: Residuals (proportions), dtype: float64

```


The model was about 66% correct on the test data.

NOTE: 0 means the prediction and the actual value matched, whereas 1 means the prediction and the actual value did not match. The accuracies are very close (ie 65% and 66%); this is a

good sign to show that there's neither underfitting nor overfitting. But, remember the churn data is imbalanced and we will later review that.

```
#Compute the confusion matrix
cm = confusion_matrix(y_test, y_pred_test)

# Display the confusion matrix
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=model.classes_)
disp.plot(cmap='Blues')
plt.title('Confusion Matrix - Logistic Regression')
plt.show()
```



We will go ahead and use cross validation to further evaluate our fitted model.

```
# Import the function
from sklearn.model_selection import cross_val_score

# Perform cross-validation
cv_scores = cross_val_score(model, X_train_new, y_train, cv=3)
print(cv_scores)
print("Mean_accuracy:", np.mean(cv_scores))

[0.66982249 0.63076923 0.63550296]
Mean_accuracy: 0.6453648915187377
```

The values are generally close which shows the model is consistent across the subsets. The mean accuracy score 64%, which is generally acceptable but we need a higher accuracy.

We can now compare the metrics values for the fitted model and the baseline metrics.

```
model_accuracy = accuracy_score(y_test, [1] * len(y_test))
model_recall = recall_score(y_test, [1] * len(y_test))
model_precision = precision_score(y_test, [1] * len(y_test))
model_f1 = f1_score(y_test, [1] * len(y_test))

print(f"""
Accuracy
Baseline: {baseline_accuracy:1.3f} Fitted Model: {model_accuracy:1.3f}
Recall
Baseline: {baseline_recall:1.3f} Fitted Model: {model_recall:1.3f}
Precision
Baseline: {baseline_precision:1.3f} Fitted Model:
{model_precision:1.3f}
F1 Score
Baseline: {baseline_f1:1.3f} Fitted Model: {model_f1:1.3f}
""")
```

```
Accuracy
Baseline: 0.845 Fitted Model: 0.155
Recall
Baseline: 1.000 Fitted Model: 1.000
Precision
Baseline: 0.845 Fitted Model: 0.155
F1 Score
Baseline: 0.916 Fitted Model: 0.268
```

The baseline metrics show high accuracy, precision and f1 scores showing generally that the model performs well. The fitted model has very low scores compared to the baseline metrics. It has a poor performance showing that it interprets most of positive values incorrectly hence the low scores in accuracy and precision.

Handling imbalance in the data.

We will use SMOTE i.e is Synthetic Minority Oversampling. We increase the number of the minority class instances.

```
from imblearn.over_sampling import SMOTE
# Original class distribution
print('Original class distribution: \n')
print(y.value_counts())
smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_new,
y_train)
# Preview synthetic sample class distribution
print('-----')
print('Synthetic sample class distribution: \n')
print(pd.Series(y_train_resampled).value_counts())

Original class distribution:

0      2727
1       442
Name: Churn, dtype: int64
-----
Synthetic sample class distribution:


1      2191
0      2191
Name: Churn, dtype: int64

logreg = LogisticRegression(fit_intercept=False, C=1e12,
solver='liblinear')
logreg.fit(X_train_resampled, y_train_resampled)
print(logreg)

# Predict
y_hat_test = logreg.predict(X_test_new)

LogisticRegression(C=10000000000000.0, fit_intercept=False,
solver='liblinear')

#Compute the confusion matrix for the resampled data
cm = confusion_matrix(y_test, y_hat_test)
disp = ConfusionMatrixDisplay(confusion_matrix=cm,
display_labels=logreg.classes_)
disp.plot(cmap='Reds')
plt.title('Confusion Matrix - Logistic Regression (SMOTE)')
plt.show()
```



```
class_report = classification_report(y_test, y_hat_test)
print(class_report)
```

	precision	recall	f1-score	support
0	0.95	0.62	0.75	536
1	0.29	0.84	0.43	98
accuracy			0.66	634
macro avg	0.62	0.73	0.59	634
weighted avg	0.85	0.66	0.70	634

We will use the ROC and AUC, to evaluate the logistic model after balancing the the dataset.

```
from sklearn.metrics import roc_curve, auc


fpr, tpr, thresholds = roc_curve(y_test, y_score)
print('AUC: {}'.format(auc(fpr, tpr)))

AUC: 0.7994593359731954

# Visualizing the ROC and AUC for Logistic Regression Smote
sns.set_style('darkgrid', {'axes.facecolor': '0.9'})

print('AUC: {}'.format(auc(fpr, tpr)))
plt.figure(figsize=(10, 8))
lw = 2
plt.plot(fpr, tpr, color='darkorange',
         lw=lw, label='ROC curve')
plt.plot([0, 1], [0, 1], color='navy', lw=lw, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.yticks([i/20.0 for i in range(21)])
plt.xticks([i/20.0 for i in range(21)])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver operating characteristic (ROC) Curve')
plt.legend(loc='lower right')
plt.show()

AUC: 0.7994593359731954
```



AUC of 0.79946 demonstrates an outstanding discriminatory power; it is capable of accurately identifying customers who will churn.

Hypertuning

Hypertuning parameters is the process of tweaking model parameters to improve its performance. It will help achieve a maximum performance. We will tweak the parameters for the SMOTE logistic regression like changing the C intercept values and specifying the L1 penalty and perform GridSearch with cross validation.

```
from sklearn.model_selection import GridSearchCV

log1 = LogisticRegression(penalty = 'l1',
                           solver='liblinear', random_state=42)
log1.fit(X_train_resampled, y_train_resampled)
param_grid = {
    'C': [0.01, 0.1, 1, 10, 100] # Different values for the inverse
    of regularization strength
}
#Perform GridSearch with cross validation.
```

```
# Set up the grid search
grid_search = GridSearchCV(logreg, param_grid, cv=5,
scoring='roc_auc', n_jobs=-1, verbose=1)

# Fit the grid search to the resampled training data
grid_search.fit(X_train_resampled, y_train_resampled)

# Get the best model
best_model = grid_search.best_estimator_

print(f'Best C parameter: {grid_search.best_params_}')
print(f'Best AUC: {grid_search.best_score_:.4f}')
```

Fitting 5 folds for each of 5 candidates, totalling 25 fits
Best C parameter: {'C': 100}
Best AUC: 0.8449

```
from sklearn.metrics import roc_auc_score
#Evaluate the model
# Make predictions on the test set
y_pred = grid_search.predict(X_test_new)
y_pred_prob = grid_search.predict_proba(X_test_new)[:, 1]
```

```
# Classification report
print("Classification Report:")
print(classification_report(y_test, y_pred))
```

```
# ROC AUC score
roc_auc = roc_auc_score(y_test, y_pred_prob)
print(f"ROC AUC: {roc_auc:.4f}")
```


```
Classification Report:
```

	precision	recall	f1-score	support
0	0.95	0.61	0.74	536
1	0.28	0.83	0.42	98
accuracy			0.65	634
macro avg	0.62	0.72	0.58	634
weighted avg	0.85	0.65	0.69	634

ROC AUC: 0.8005

```
fpr, tpr, _ = roc_curve(y_test, y_pred_prob)
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (AUC = {roc_auc:.4f})')
plt.plot([0, 1], [0, 1], color='grey', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



After tuning the parameters, we see the training set , after using different C intercept values we get the best parameter combination AUC of 0.8455. On the test set we get an AUC of 0.8004. The slight difference on the AUC shows there's no overfitting. 0.8455 shows that the model has an impressive ability to distinguish between the churn and no churn within the training dataset. An AUC of 0.8004 on the test set shows that the model generalizes well to the unseen data hence validating the model's use for practical use. It definitely performs better than the original logistic regression mode with the AUC of 0.799945

```
from sklearn.inspection import permutation_importance
perm_importance = permutation_importance(log1, X_train_resampled,
y_train_resampled, n_repeats=10, random_state=42)

# Extract importance scores
importance_scores = perm_importance.importances_mean

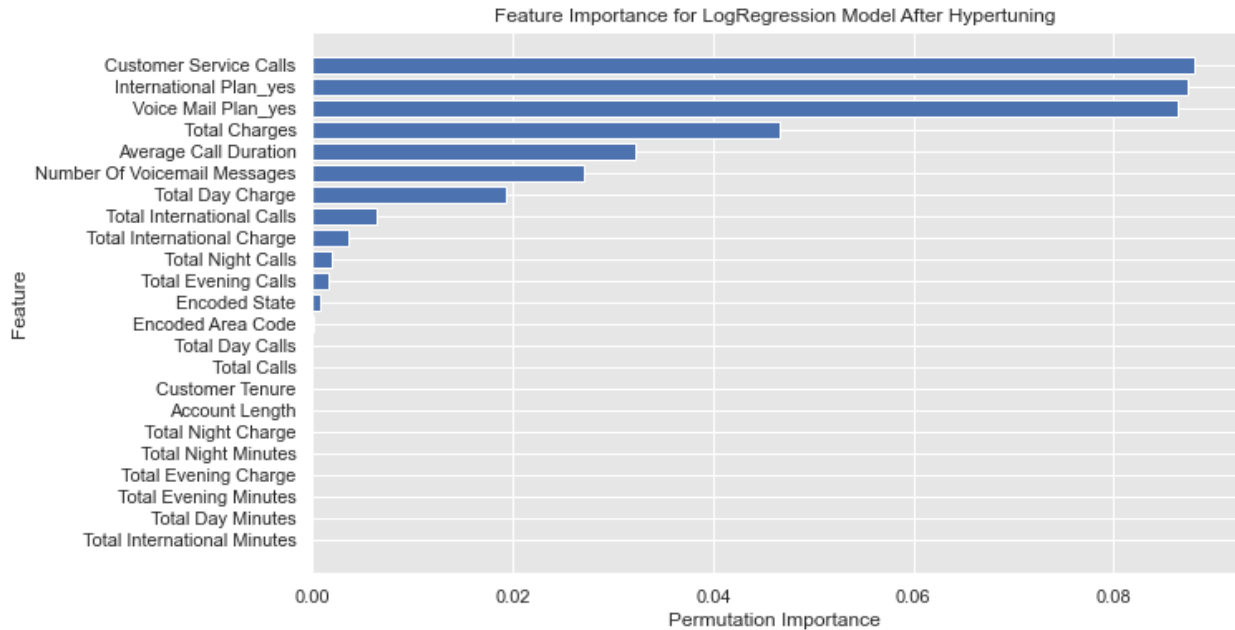
# Create a DataFrame for better visualization
feature_importance_df = pd.DataFrame({
    'Feature': X_train_resampled.columns,
    'Importance': importance_scores
}).sort_values(by='Importance', ascending=False)
```

```
print(feature_importance_df)
```

	Feature	Importance
14	Customer Service Calls	0.087951
22	International Plan_yes	0.087449
21	Voice Mail Plan_yes	0.086353
15	Total Charges	0.046737
17	Average Call Duration	0.032337
1	Number Of Voicemail Messages	0.027134
4	Total Day Charge	0.019283
12	Total International Calls	0.006458
13	Total International Charge	0.003651
9	Total Night Calls	0.001963
6	Total Evening Calls	0.001575
19	Encoded State	0.000799
20	Encoded Area Code	0.000183
3	Total Day Calls	0.000068
16	Total Calls	0.000000
18	Customer Tenure	0.000000
0	Account Length	0.000000
10	Total Night Charge	0.000000
8	Total Night Minutes	0.000000
7	Total Evening Charge	0.000000
5	Total Evening Minutes	0.000000
2	Total Day Minutes	0.000000
11	Total International Minutes	0.000000

```
# Plot the feature importance
```

```
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df['Feature'],
feature_importance_df['Importance'])
plt.xlabel('Permutation Importance')
plt.ylabel('Feature')
plt.title('Feature Importance for LogRegression Model After
Hypertuning')
plt.gca().invert_yaxis()
plt.show()
```

K-Nearest Neighbours.

This is a more complex model. It relies on the idea that similar data points tend to have similar labels or values. We will use it to make more predictions and compare the results to the logistic regression.

```
from sklearn.pipeline import Pipeline
from sklearn.neighbors import KNeighborsClassifier
#Initialize the KNN model
knn = KNeighborsClassifier(n_neighbors=5)

# Train the model on the resampled training set
knn.fit(X_train_resampled, y_train_resampled)

# Define the pipeline
pipeline = Pipeline([
    ('knn', KNeighborsClassifier())
])

# Define the hyperparameter grid
param_grid = {
    'knn__n_neighbors': [3, 5, 7, 9],
    'knn__weights': ['uniform', 'distance'],
    'knn__metric': ['euclidean', 'manhattan']
}

# Initialize GridSearchCV
grid_search = GridSearchCV(pipeline, param_grid, cv=5,
    scoring='roc_auc', n_jobs=-1, verbose=2)

# Perform the grid search
```

```

grid_search.fit(X_train_resampled, y_train_resampled)

# Get the best parameters and best score
best_params = grid_search.best_params_
best_score = grid_search.best_score_
print(f"Best parameters: {best_params}")
print(f"Best ROC AUC: {best_score}")

Fitting 5 folds for each of 16 candidates, totalling 80 fits
Best parameters: {'knn__metric': 'manhattan', 'knn__n_neighbors': 9,
'knn__weights': 'distance'}
Best ROC AUC: 0.9882564713385434

# Train the final model with the best parameters on the entire
training set
final_model = grid_search.best_estimator_

# Evaluate the final model
# Make predictions on the training set
y_train_pred = final_model.predict(X_train_resampled)
y_train_pred_proba = final_model.predict_proba(X_train_resampled)[: ,
1]

# Make predictions on the test set
y_test_pred = final_model.predict(X_test_new)
y_test_pred_proba = final_model.predict_proba(X_test_new)[: , 1]

# Generate classification report for training set
print("Classification Report for Training Set:")
print(classification_report(y_train_resampled, y_train_pred))

# Generate classification report for test set
print("Classification Report for Test Set:")
print(classification_report(y_test, y_test_pred))

```

Classification Report for Training Set:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	2191
1	1.00	1.00	1.00	2191
accuracy			1.00	4382
macro avg	1.00	1.00	1.00	4382
weighted avg	1.00	1.00	1.00	4382

Classification Report for Test Set:

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.94	0.82	0.88	536
1	0.42	0.69	0.52	98
accuracy			0.80	634
macro avg	0.68	0.76	0.70	634
weighted avg	0.86	0.80	0.82	634

#ROC AUC for training set

```
roc_auc_train = roc_auc_score(y_train_resampled, y_train_pred_proba)
print(f"AUC for Training Set: {roc_auc_train:.4f}")
```

ROC AUC score for test set


```
roc_auc_test = roc_auc_score(y_test, y_test_pred_proba)
print(f"AUC for Test Set: {roc_auc_test:.4f}")
```

AUC for Training Set: 1.0000

AUC for Test Set: 0.8256

Plot the confusion matrix for the test set


```
ConfusionMatrixDisplay.from_predictions(y_test, y_test_pred)
plt.title('Confusion Matrix for Test Set')
plt.show()
```



```
# Plot the ROC curve
from sklearn.metrics import roc_curve

fpr, tpr, _ = roc_curve(y_test, y_test_pred_proba)
plt.figure(figsize=(10, 6))
plt.plot(fpr, tpr, color='blue', lw=2, label=f'ROC curve (AUC = {roc_auc_test:.4f})')
plt.plot([0, 1], [0, 1], color='grey', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
```

```
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic (ROC) Curve')
plt.legend(loc="lower right")
plt.show()
```



```
from sklearn.inspection import permutation_importance
perm_importance = permutation_importance(knn, X_train_resampled,
y_train_resampled, n_repeats=10, random_state=42)

# Extract importance scores
importance_scores = perm_importance.importances_mean

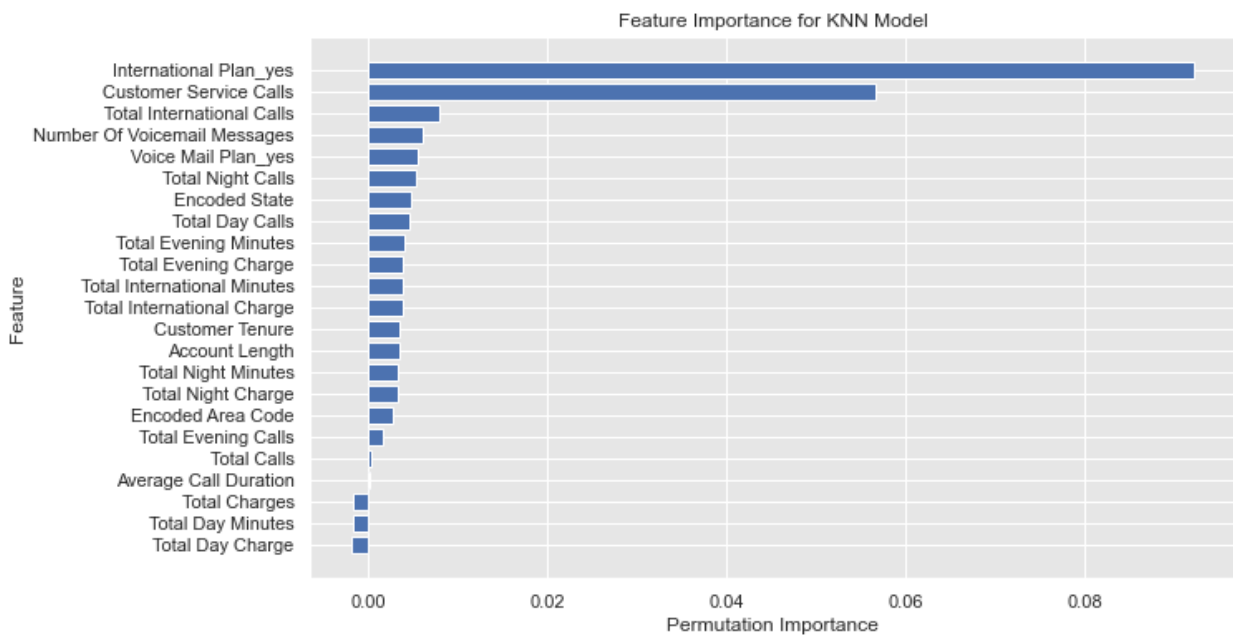
# Create a DataFrame for better visualization
feature_importance_df = pd.DataFrame({
    'Feature': X_train_resampled.columns,
    'Importance': importance_scores
}).sort_values(by='Importance', ascending=False)

print(feature_importance_df)
```

	Feature	Importance
22	International Plan_yes	0.092173
14	Customer Service Calls	0.056709
12	Total International Calls	0.007964
1	Number Of Voicemail Messages	0.006162
21	Voice Mail Plan_yes	0.005568

9	Total Night Calls	0.005431
19	Encoded State	0.004792
3	Total Day Calls	0.004564
5	Total Evening Minutes	0.004016
7	Total Evening Charge	0.003994
11	Total International Minutes	0.003971
13	Total International Charge	0.003857
18	Customer Tenure	0.003583
0	Account Length	0.003583
8	Total Night Minutes	0.003400
10	Total Night Charge	0.003355
20	Encoded Area Code	0.002875
6	Total Evening Calls	0.001712
16	Total Calls	0.000411
17	Average Call Duration	0.000137
15	Total Charges	-0.001597
2	Total Day Minutes	-0.001712
4	Total Day Charge	-0.001757

```
# Plot the feature importance
plt.figure(figsize=(10, 6))
plt.barh(feature_importance_df['Feature'],
feature_importance_df['Importance'])
plt.xlabel('Permutation Importance')
plt.ylabel('Feature')
plt.title('Feature Importance for KNN Model')
plt.gca().invert_yaxis()
plt.show()
```



Summary

For all the models we have used, the results show low scores on the precisions with the highest being 0.42 of the KNN. This means that it can only 42% of the predicted values were correctly classified. It has the highest f1 value of 0.52, while the model with the second highest F1 score is the logistic regression(SMOTE) with a score of 0.43. This shows a low imbalance between precision and recall scores.

The model with the highest recall value is logistic(SMOTE) with a score of 0.84. This means that it correctly identified, captured the actual churners in the dataset but at the expense of precision as it has a low precision of 0.29. The classification metrics between the logistic(SMOTE) and logistic(SMOTE) after hypertuning has a slight difference of 0.1.

Recall is more useful for this analysis as we want to actually get the actual positive values. The models best for predicting unseen data are the LogisticRegression(SMOTE) and LogisticRegression(SMOTE) after hypertuning.

If we consider the ROC AUC values, the Logistic Regression after hypertuning showed 0.8449 on the best parameter combination for the training set and 0.80044 on the test set. There is a slight difference and overfitting should not be a big deal. Compared to the KNN, the score was 0.9882 on the best parameter combination on the train set. This showed that it performed really well on the training set. However, as much as it has a 0.8256 AUC score on the test set, which is still good, there is a significant difference between the two scores signifying overfitting.

With these results, the best models to be used to predict the unseen data for churners is the Logistic Regression(SMOTE) and LogisticRegression(SMOTE) after hypertuning.

Conclusions

In Conclusion, our goal was to identify clients which are likely to churn, so we can do special-purpose marketing strategies to avoid the churn event. From the feature importance graphs shown for both of the models, we can see which features contribute most to the prediction of the models. The top 4 are Customer Service Calls, International Plan, Voice Mail Plan and Total International Calls. SyriaTel can improve their products by focusing on these features.

Limitations

1. Overfitting - The KNN performed well on the train data, but reduced its generalization on the test data.
2. GridSearch with cross validation is time consuming hence could not build more models.
3. Imbalanced Data- SMOTE might not perfectly represent the true distribution.
4. Feature Selection method did not really have an effect to the data. It did not reduce the relevant features from the dataset.

Recommendations.

1. Use a more updated dataset to reduce the imbalance of data.
2. Use better and more convenient method for feature selection to reducing overfitting.